

# MA 5124 Financial Time Series Analysis and Forecasting

## Chapter 1: Introduction to time series and forecasting Lesson 2

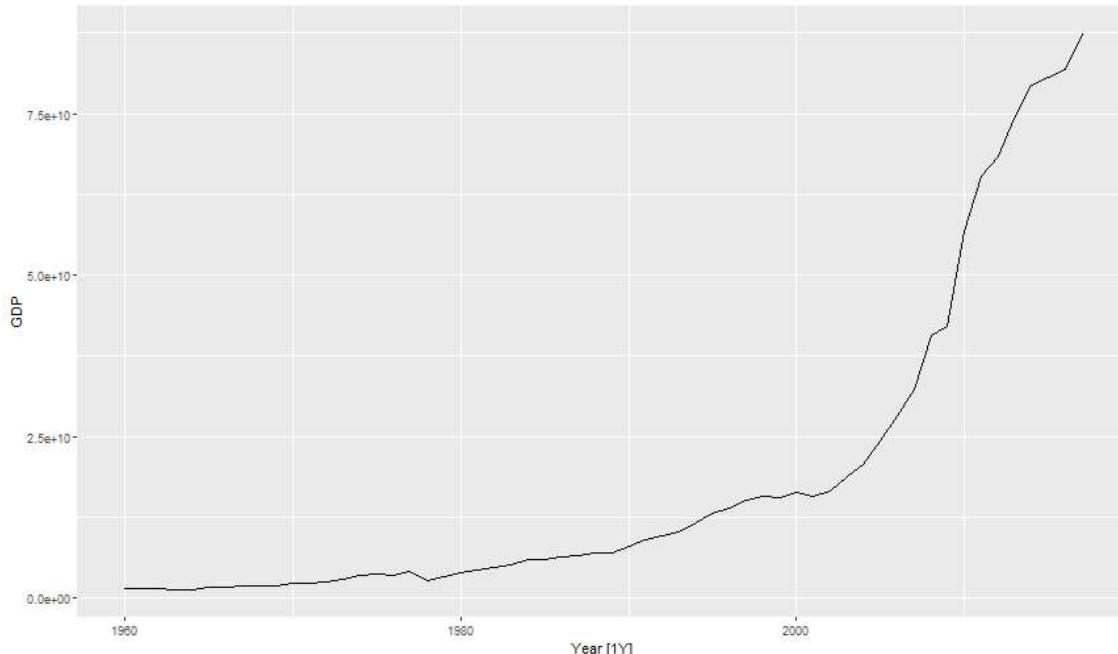
Dr. Priyanga Talagala

30-06-2024

# Transformations and adjustments

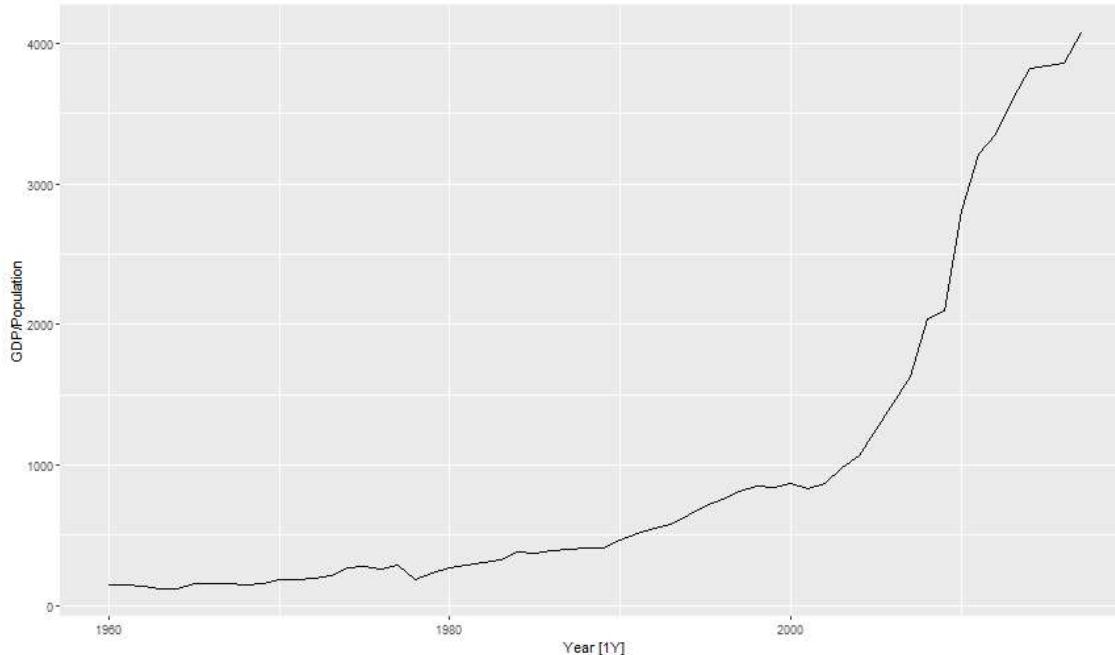
# Per capita adjustments

```
global_economy |>
  filter(Country == "Sri Lanka") |>
  autoplot(GDP)
```



# Per capita adjustments

```
global_economy |>  
  filter(Country == "Sri Lanka") |>  
  autoplot(GDP / Population)
```



# Mathematical transformations

- ▶ If the data show different variation at different levels of the series, then a transformation can be useful.
- ▶ Denote original observations as  $y_1, \dots, y_n$  and transformed observations as  $w_1, \dots, w_n$ .
- ▶ Mathematical transformations for stabilizing variation

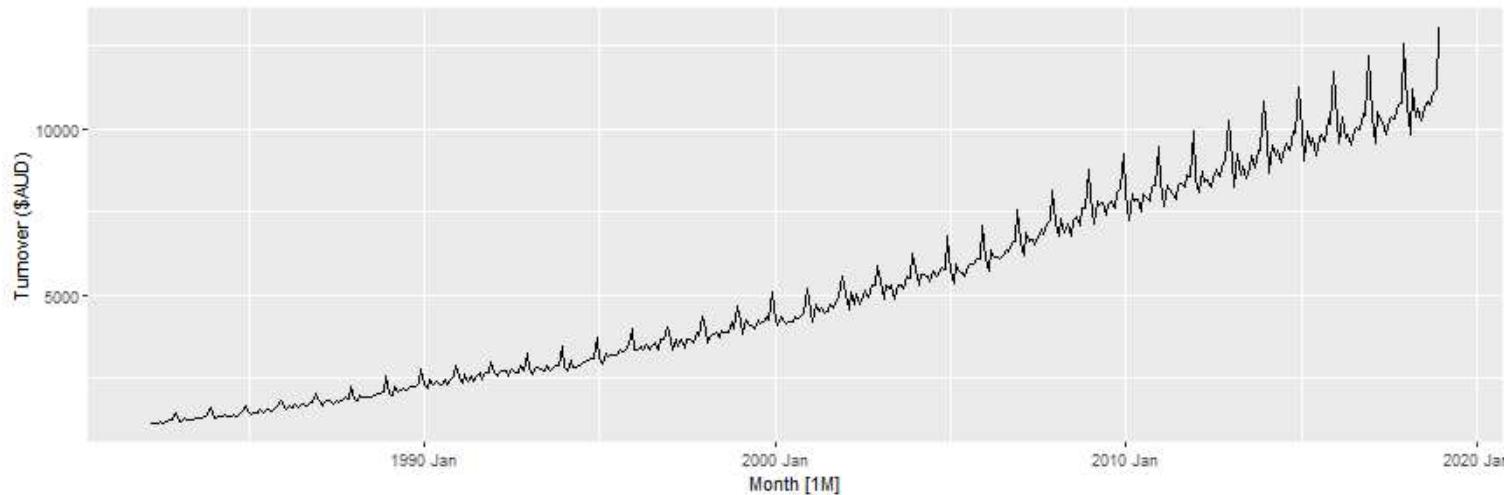
Transformation	Equation
Square root	$w_t = \sqrt{y_t}$
Cube root	$w_t = \sqrt[3]{y_t}$
Logarithm	$w_t = \log(y_t)$

- ▶ Logarithms, in particular, are useful because they are more interpretable: changes in a log value are **relative (percent) changes on the original scale**.

# Mathematical transformations

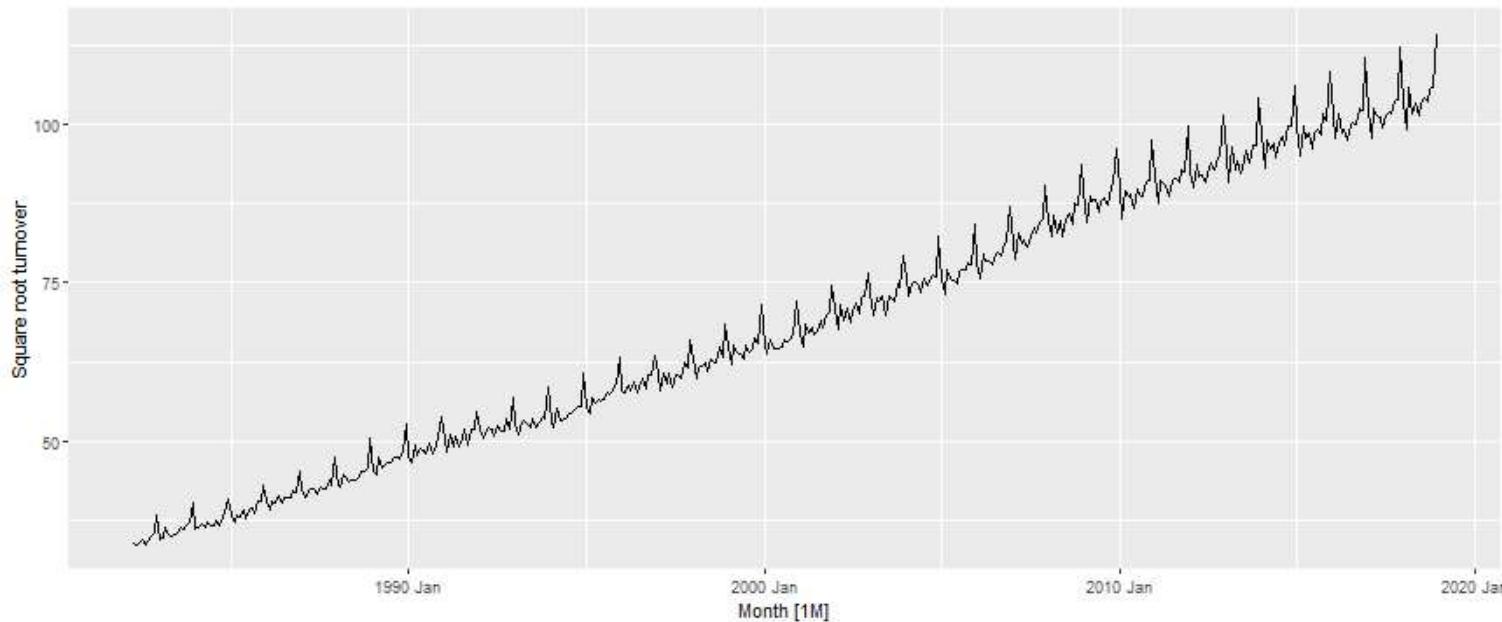
```
food <- aus_retail |>
  filter(Industry == "Food retailing") |>
  summarise(Turnover = sum(Turnover))
```

```
food |> autoplot(Turnover) +
  labs(y = "Turnover ($AUD)")
```



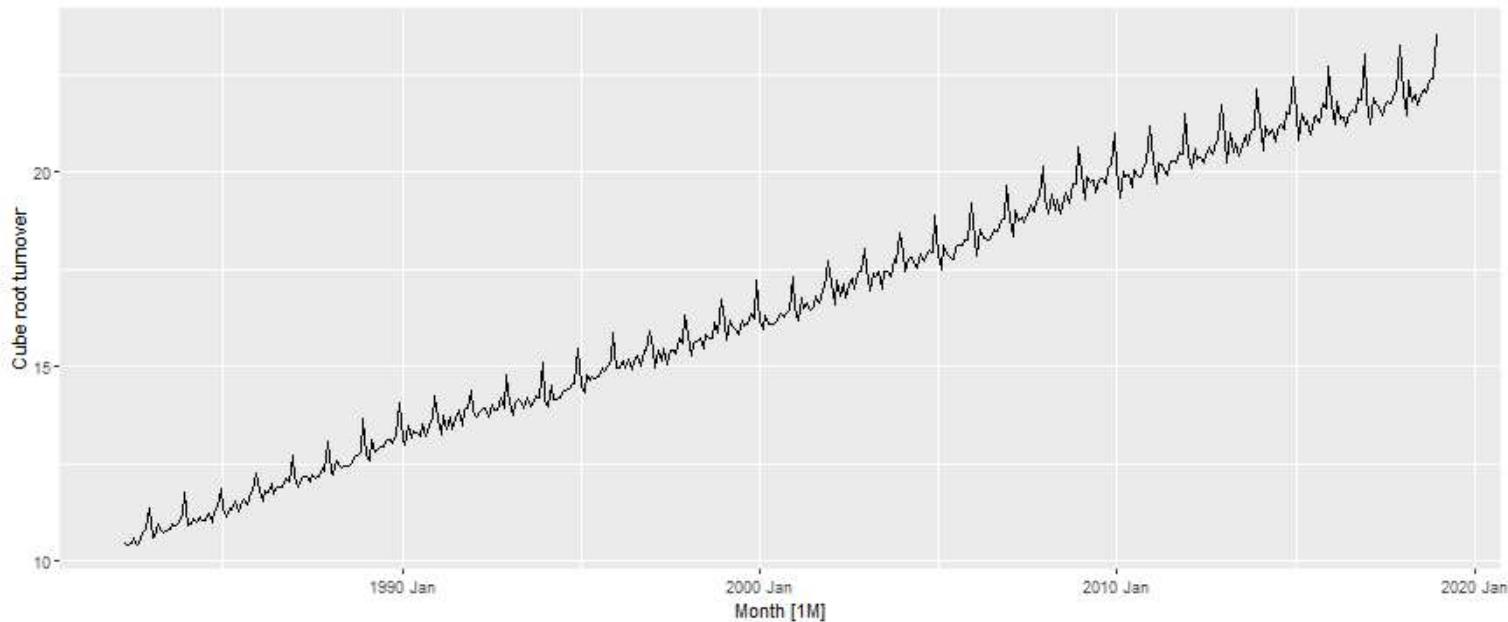
# Mathematical transformations

```
food |> autoplot(sqrt(Turnover)) +  
  labs(y = "Square root turnover")
```



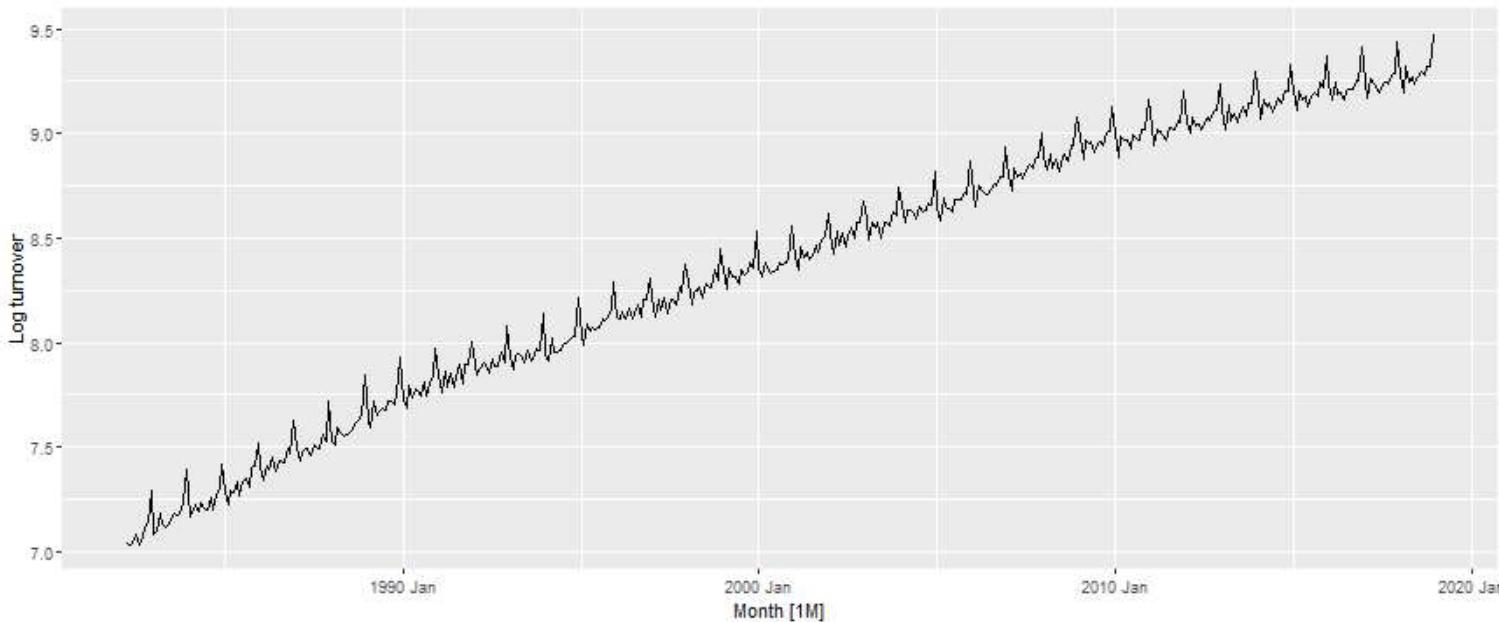
# Mathematical transformations

```
food |> autoplot(Turnover^(1/3)) +  
  labs(y = "Cube root turnover")
```



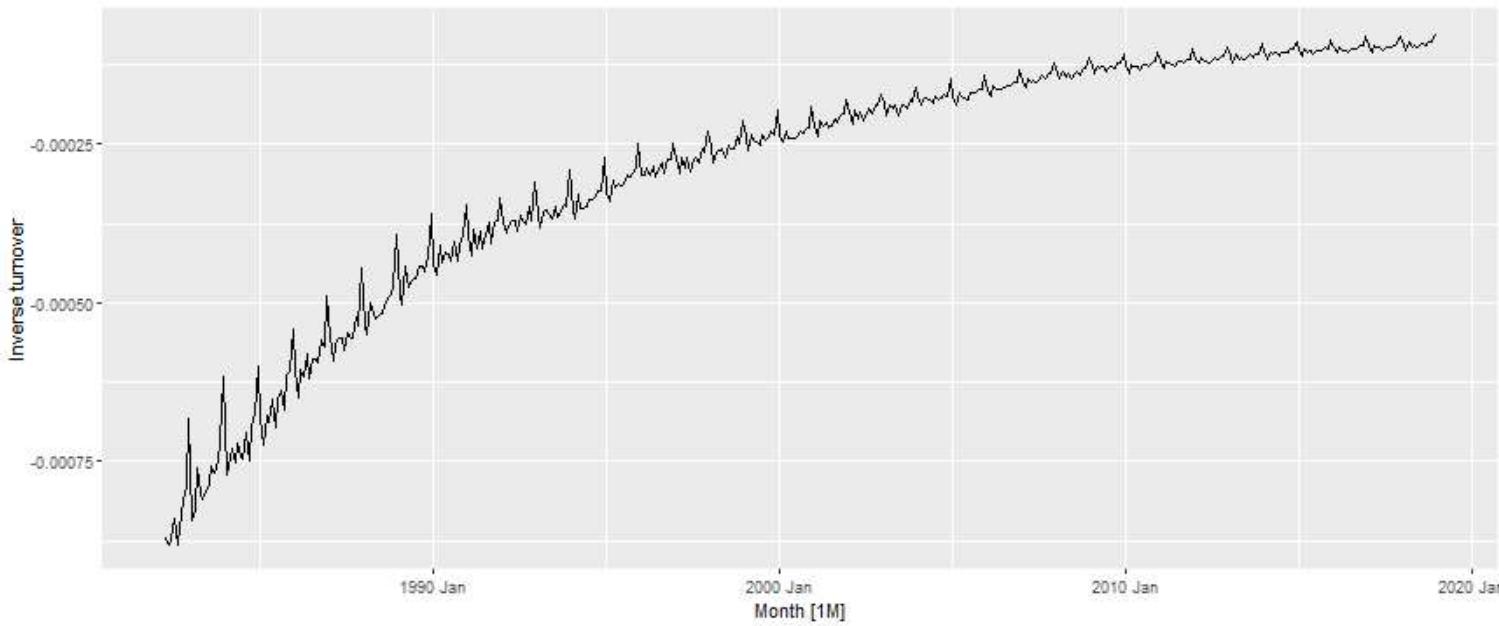
# Mathematical transformations

```
food |> autoplot(log(Turnover)) +  
  labs(y = "Log turnover")
```



# Mathematical transformations

```
food |> autoplot(-1/Turnover) +  
  labs(y = "Inverse turnover")
```



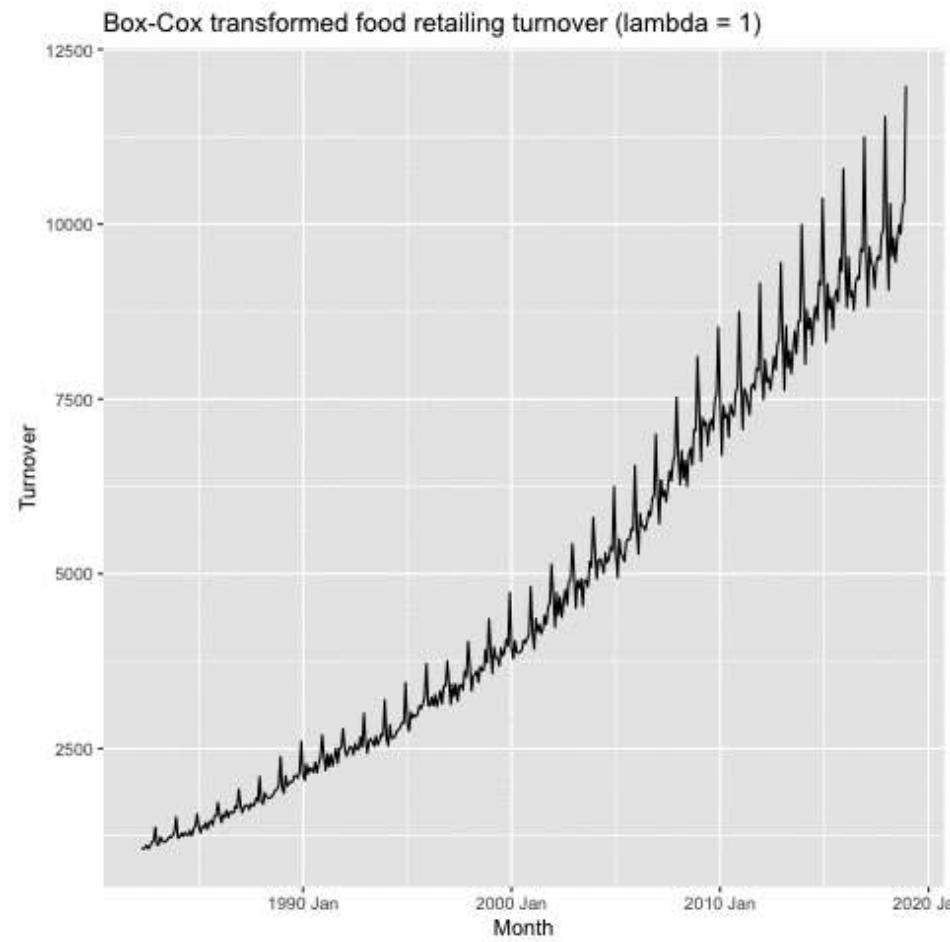
# Box-Cox transformations

- ▶ Each of these transformations is close to a member of the family of **Box-Cox transformations**:

$$w_t = \begin{cases} \log(y_t), & \lambda = 0; \\ (y_t^\lambda - 1)/\lambda, & \lambda \neq 0. \end{cases}$$

- ▶  $\lambda = 1$ : (No substantive transformation)
- ▶  $\lambda = \frac{1}{2}$ : (Square root plus linear transformation)
- ▶  $\lambda = 0$ : (Natural logarithm)
- ▶  $\lambda = -1$ : (Inverse plus 1)

# Box-Cox transformations



# Box-Cox transformations

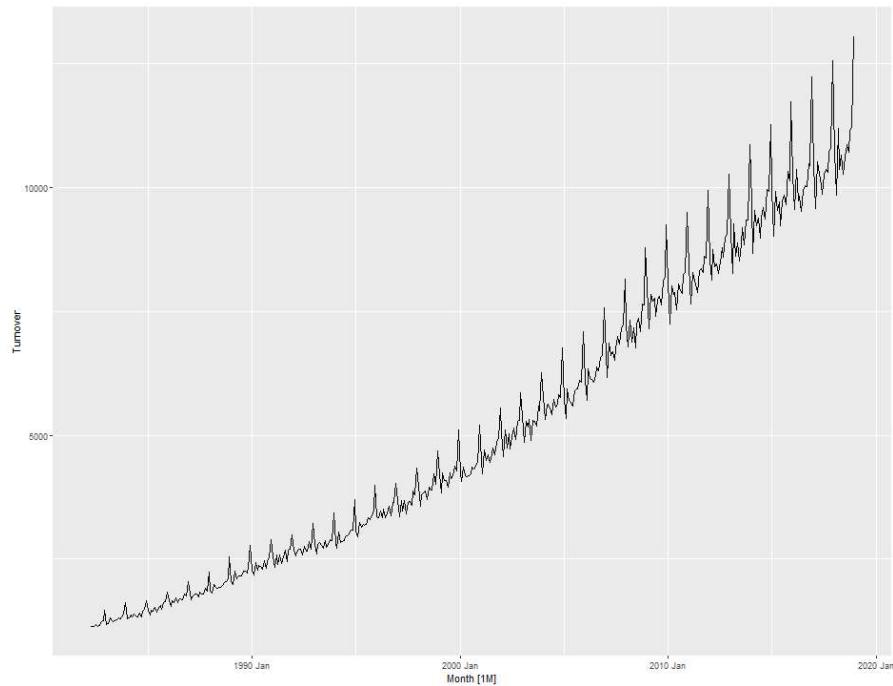
```
food |>  
  features(Turnover, features = guererro)
```

```
## # A tibble: 1 × 1  
##   lambda_guererro  
##       <dbl>  
## 1      0.0895
```

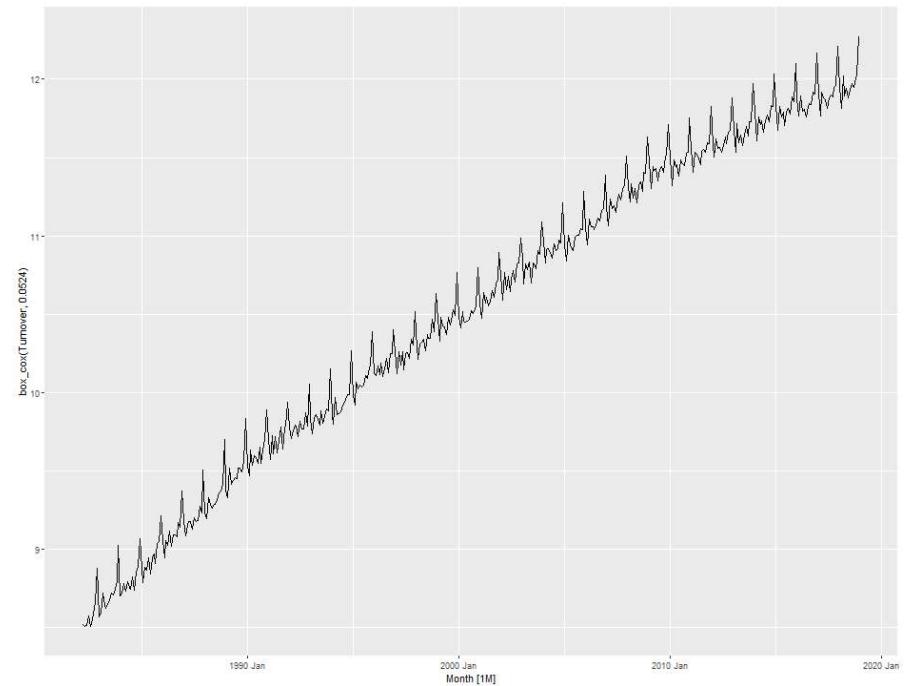
- ▶ This attempts to balance the seasonal fluctuations and random variation across the series.
- ▶ Always check the results.
- ▶ A low value of  $\lambda$  can give extremely large prediction intervals.

# Box-Cox transformations

```
food |>  
  autoplot(Turnover)
```



```
food |>  
  autoplot(box_cox(Turnover, 0.0524))
```



# Transformations

- ▶ Often no transformation needed.
- ▶ Simple transformations are easier to explain and work well enough.
- ▶ Transformations can have very large effect on PI.
- ▶ If the data contains zeros, then don't take logs.
- ▶ `log1p()` can be useful for data with zeros.
- ▶ If some data are negative, no power transformation is possible unless a constant is added to all values.
- ▶ Choosing logs is a simple way to force forecasts to be positive
- ▶ Transformations must be reversed to obtain forecasts on the original scale. (Handled automatically by `fable`.)

# Time series components

## RECALL: Time series patterns

- ▶ **Trend** pattern exists when there is a long-term increase or decrease in the data.
- ▶ **Seasonal** pattern exists when a series is influenced by seasonal factors (e.g., the quarter of the year, the month, or day of the week).
- ▶ **Cyclic** pattern exists when data exhibit rises and falls that are not of fixed frequency (duration usually of at least 2 years).

# Time series decomposition

$$y_t = f(S_t, T_t, R_t)$$

where

- ▶  $y_t$  = data at period  $t$
- ▶  $T_t$  = trend-cycle component at period  $t$
- ▶  $S_t$  = seasonal component at period  $t$
- ▶  $R_t$  = remainder component at period  $t$

**Additive decomposition:**  $y_t = S_t + T_t + R_t.$

**Multiplicative decomposition:**  $y_t = S_t \times T_t \times R_t.$

# Time series decomposition

- ▶ Additive model appropriate if magnitude of seasonal fluctuations does not vary with level.
- ▶ If seasonal are proportional to level of series, then multiplicative model appropriate.
- ▶ Multiplicative decomposition more prevalent with economic series
- ▶ Alternative: use a Box-Cox transformation, and then use additive decomposition.
- ▶ Logs turn multiplicative relationship into an additive relationship:

$$y_t = S_t \times T_t \times E_t \quad \Rightarrow \quad \log y_t = \log S_t + \log T_t + \log E_t.$$

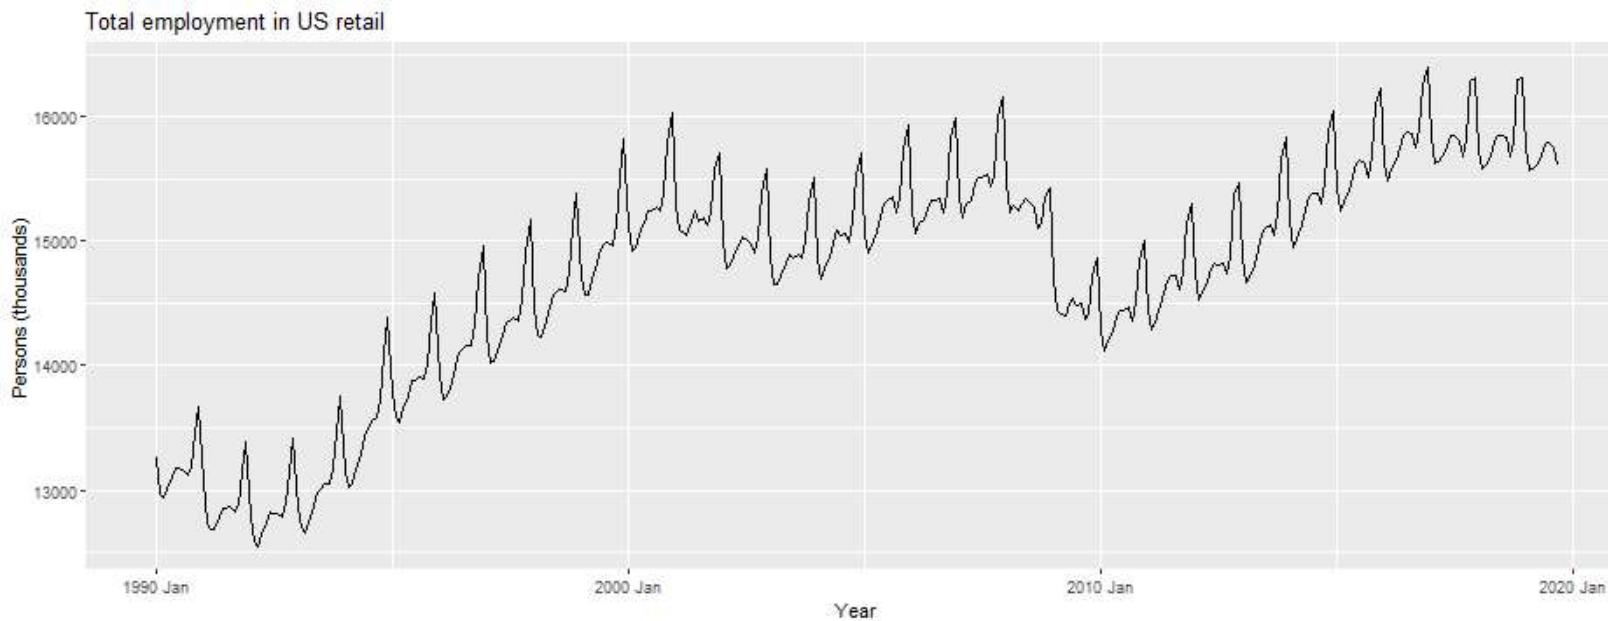
# US Retail Employment

```
library(fpp3)
us_retail_employment <- us_employment |>
  filter(year(Month) >= 1990, Title == "Retail Trade") |>
  select(-Series_ID)
us_retail_employment

## # A tsibble: 357 x 3 [1M]
##       Month Title     Employed
##       <mth> <chr>     <dbl>
## 1 1990 Jan Retail Trade 13256.
## 2 1990 Feb Retail Trade 12966.
## 3 1990 Mar Retail Trade 12938.
## 4 1990 Apr Retail Trade 13012.
## 5 1990 May Retail Trade 13108.
## 6 1990 Jun Retail Trade 13183.
## 7 1990 Jul Retail Trade 13170.
## 8 1990 Aug Retail Trade 13160.
## 9 1990 Sep Retail Trade 13113.
## 10 1990 Oct Retail Trade 13185.
## # i 347 more rows
```

# US Retail Employment

```
us_retail_employment |>
  autoplot(Employed) +
  xlab("Year") + ylab("Persons (thousands)") +
  ggtitle("Total employment in US retail")
```



# US Retail Employment

```
us_retail_employment |>  
  model(stl = STL(Employed))
```

```
## # A mable: 1 x 1  
##       stl  
##     <model>  
## 1     <STL>
```

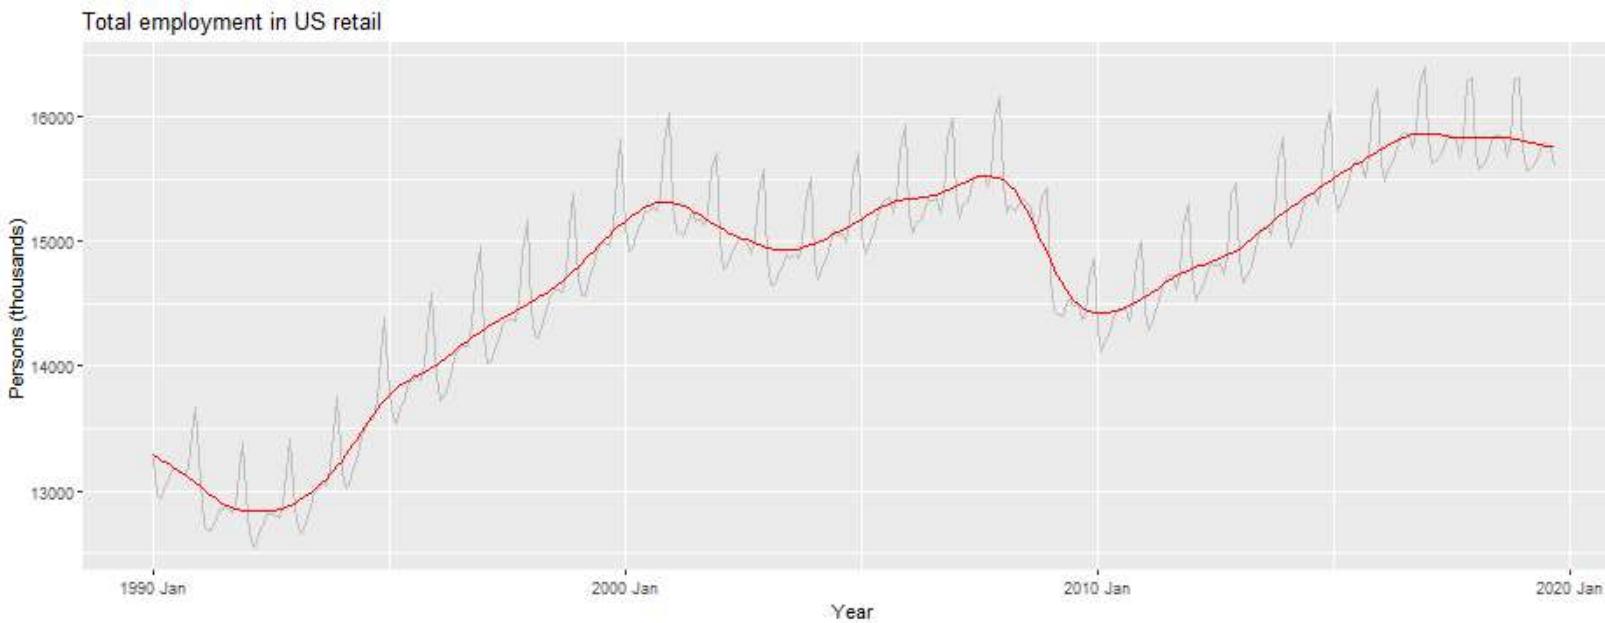
# US Retail Employment

```
dcmp <- us_retail_employment |>
  model(stl = STL(Employed))
components(dcmp)

## # A dable: 357 x 7 [1M]
## # Key:   .model [1]
## # :     Employed = trend + season_year + remainder
## #   .model Month Employed trend season_year remainder season_adjust
## #   <chr>   <mth>   <dbl>   <dbl>   <dbl>   <dbl>
## 1 stl    1990 Jan    13256. 13288.   -33.0    0.836  13289.
## 2 stl    1990 Feb    12966. 13269.   -258.    -44.6  13224.
## 3 stl    1990 Mar    12938. 13250.   -290.    -22.1  13228.
## 4 stl    1990 Apr    13012. 13231.   -220.     1.05  13232.
## 5 stl    1990 May    13108. 13211.   -114.    11.3   13223.
## 6 stl    1990 Jun    13183. 13192.   -24.3    15.5  13207.
## 7 stl    1990 Jul    13170. 13172.   -23.2    21.6   13193.
## 8 stl    1990 Aug    13160. 13151.   -9.52   17.8   13169.
## 9 stl    1990 Sep    13113. 13131.   -39.5   22.0   13153.
## 10 stl   1990 Oct    13185. 13110.   61.6    13.2   13124.
## # i 347 more rows
```

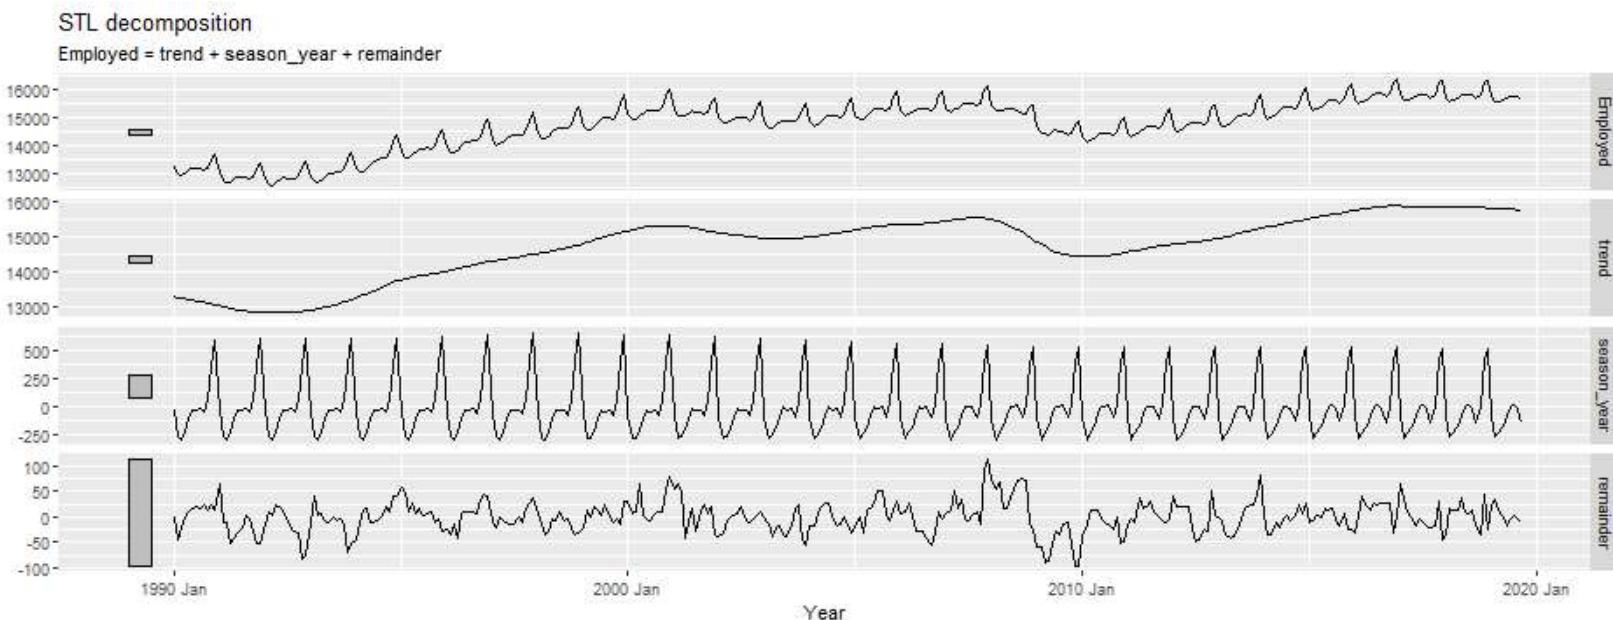
# US Retail Employment

```
us_retail_employment |>  
  autoplot(Employed, color='gray') +  
  autolayer(components(dcmp), trend, color='red') +  
  xlab("Year") + ylab("Persons (thousands)") +  
  ggtitle("Total employment in US retail")
```



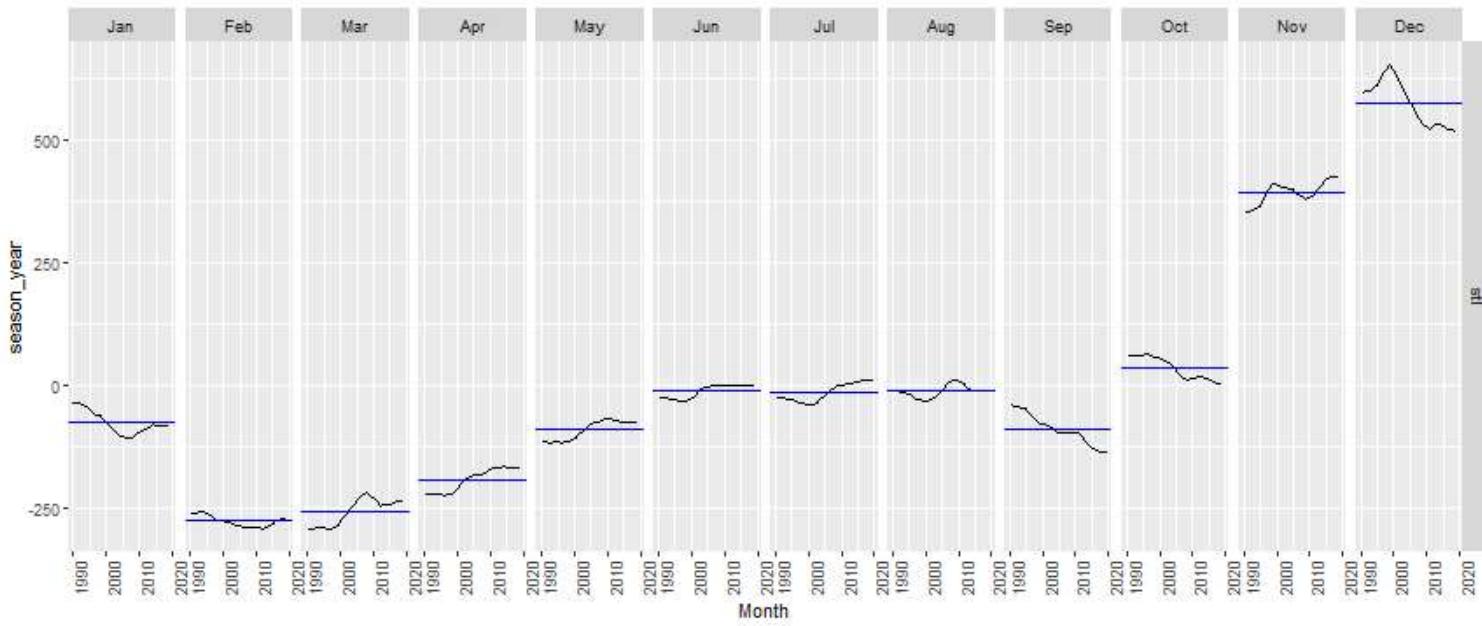
# US Retail Employment

```
components(dcmp) |> autoplot() + xlab("Year")
```



# US Retail Employment

```
components(dcmp) |> gg_subseries(season_year)
```



# Seasonal adjustment

- ▶ Useful by-product of decomposition: an easy way to calculate seasonally adjusted data.
- ▶ Additive decomposition: seasonally adjusted data given by

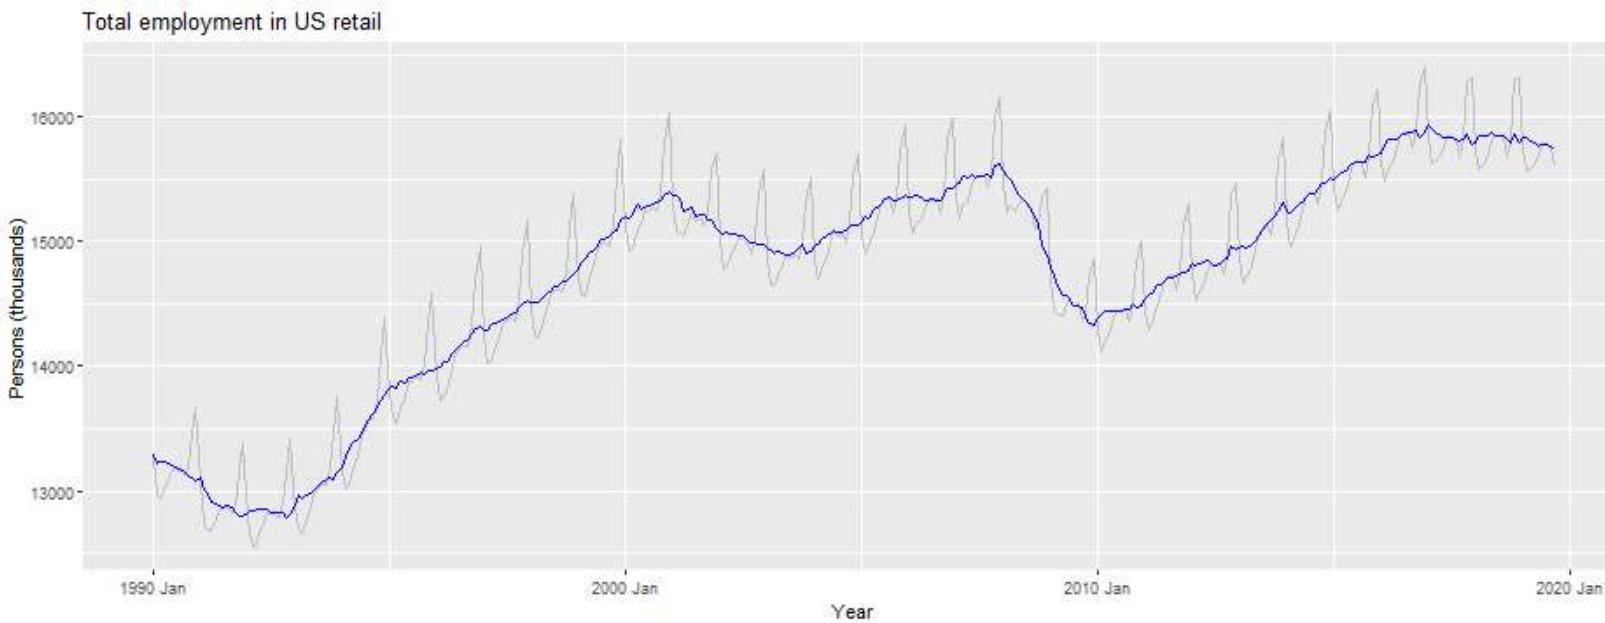
$$y_t - S_t = T_t + R_t$$

- ▶ Multiplicative decomposition: seasonally adjusted data given by

$$y_t / S_t = T_t \times R_t$$

# US Retail Employment

```
us_retail_employment |>  
  autoplot(Employed, color='gray') +  
  autolayer(components(dcmp), season_adjust, color='blue') +  
  xlab("Year") + ylab("Persons (thousands)") +  
  ggtitle("Total employment in US retail")
```



# Seasonal adjustment

- ▶ We use estimates of  $S$  based on past values to seasonally adjust a current value.
- ▶ Seasonally adjusted series reflect **remainders** as well as **trend**. Therefore they are not "smooth" and "downturns" or "upturns" can be misleading.
- ▶ It is better to use the trend-cycle component to look for turning points.

# Moving Averages

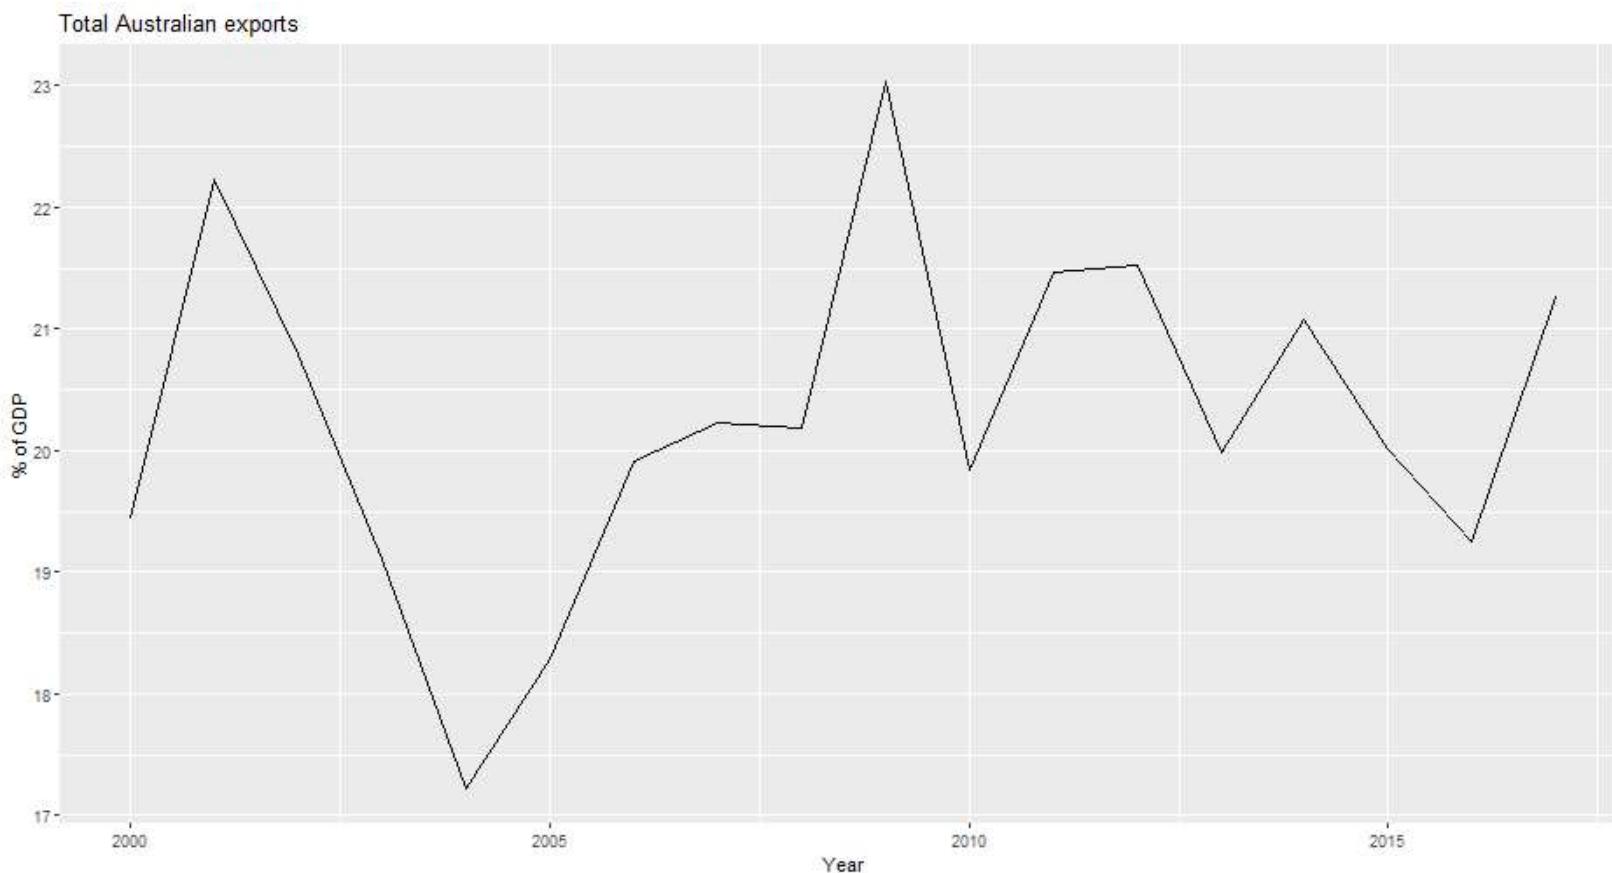
- The simplest estimate of the trend-cycle uses **moving averages**.

*m*-MA

$$\hat{T}_t = \frac{1}{m} \sum_{j=-k}^k y_{t+j}$$

where  $m = 2k + 1$

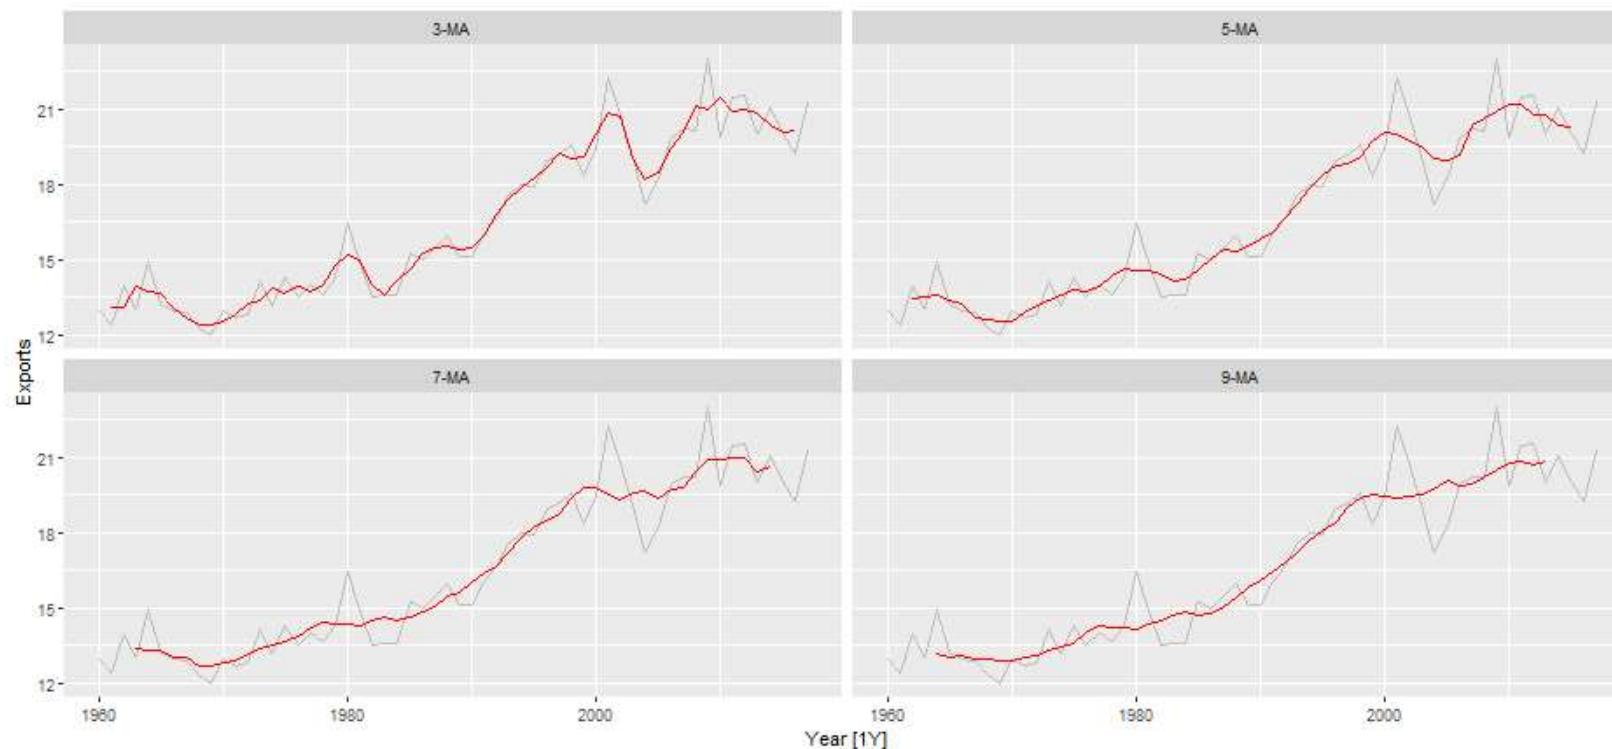
# Moving averages



# Moving averages

```
## # A tsibble: 18 x 5 [1Y]
##   Year Exports `3-MA` `5-MA` `7-MA`
##   <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1 2000     19.4     NA      NA      NA
## 2 2001     22.2     20.8     NA      NA
## 3 2002     20.8     20.7     19.8     NA
## 4 2003     19.1     19.0     19.5     19.6
## 5 2004     17.2     18.2     19.1     19.7
## 6 2005     18.3     18.5     18.9     19.4
## 7 2006     19.9     19.5     19.2     19.7
## 8 2007     20.2     20.1     20.3     19.8
## 9 2008     20.2     21.2     20.6     20.4
## 10 2009    23.0     21.0     21.0     20.9
## 11 2010    19.8     21.5     21.2     20.9
## 12 2011    21.5     20.9     21.2     21.0
## 13 2012    21.5     21.0     20.8     21.0
## 14 2013    20.0     20.9     20.8     20.5
## 15 2014    21.1     20.4     20.4     20.7
## 16 2015    20.0     20.1     20.3     NA
## 17 2016    19.3     20.2     NA      NA
## 18 2017    21.3     NA      NA      NA
```

# Moving averages



# Moving averages

- ▶ So a moving average is an **average of nearby points**
- ▶ observations nearby in time are also likely to be **close in value**.
- ▶ average eliminates some **randomness** in the data, leaving a **smooth trend-cycle** component.

$$3 - MA : \hat{T}_t = \frac{(y_{t-1} + y_t + y_{t+1})}{3}$$

$$5 - MA : \hat{T}_t = \frac{(y_{t-2} + y_{t-1} + y_t + y_{t+1} + y_{t+2})}{5}$$

- ▶ each average computed by dropping **oldest** observation and including **next** observation.
- ▶ averaging **moves** through time series until trend-cycle computed at each observation possible

# Endpoints

## Why is there no estimate at ends

- ▶ For a 3-MA, there cannot be estimates at time 1 or time  $n$  because the observations at time 0 and  $n + 1$  are not available.
- ▶ Generally: there cannot be estimates at times near the endpoints.

## The order of the MA

- ▶ larger order means smoother, flatter curve
- ▶ larger order means more points lost at ends

# Moving averages of moving averages

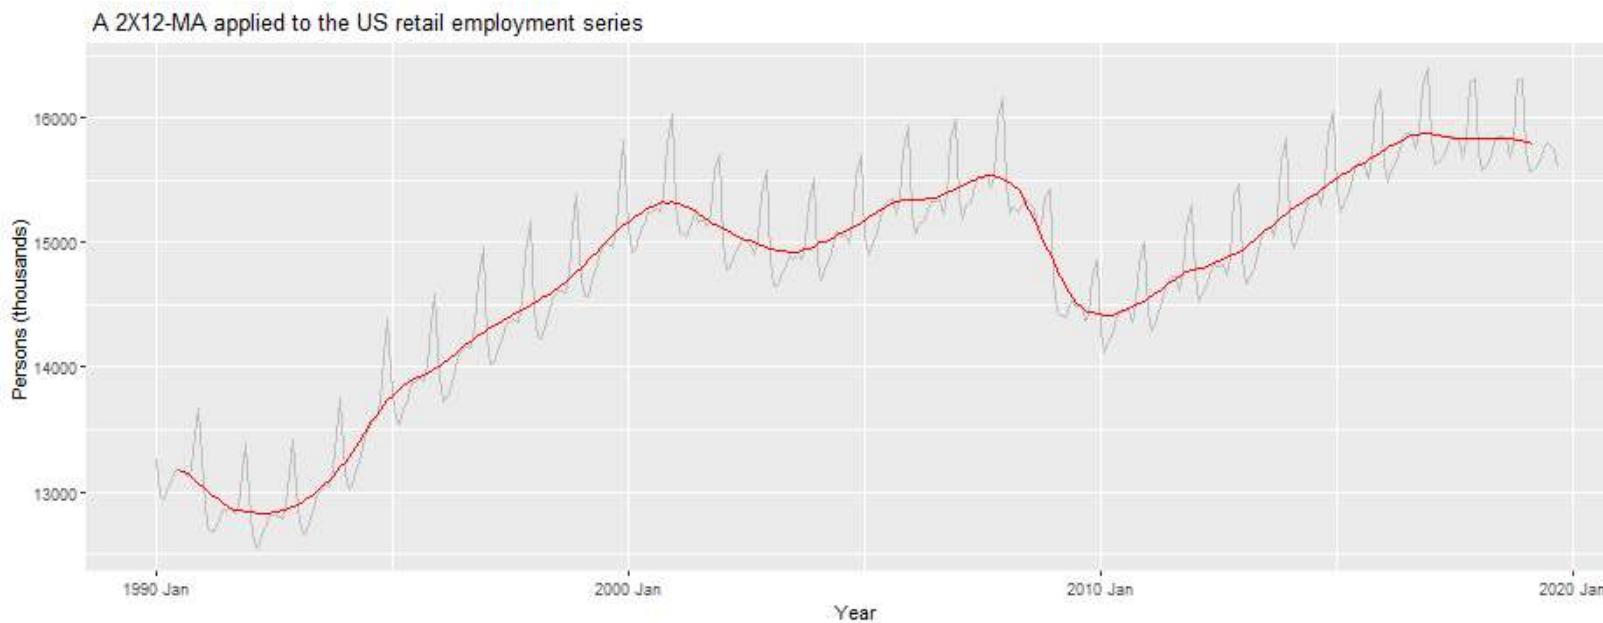
- Centered MA

## Moving averages

```
## # A tsibble: 18 x 4 [1Y]
##   Year Exports `4-MA` `2*4-MA`
##   <dbl>    <dbl>    <dbl>    <dbl>
## 1 2000     19.4     NA      NA
## 2 2001     22.2     20.4     NA
## 3 2002     20.8     19.8     20.1
## 4 2003     19.1     18.8     19.3
## 5 2004     17.2     18.6     18.7
## 6 2005     18.3     18.9     18.8
## 7 2006     19.9     19.7     19.3
## 8 2007     20.2     20.8     20.2
## 9 2008     20.2     20.8     20.8
## 10 2009    23.0     21.1     21.0
## 11 2010    19.8     21.5     21.3
## 12 2011    21.5     20.7     21.1
## 13 2012    21.5     21.0     20.9
## 14 2013    20.0     20.6     20.8
```

# Estimating the trend-cycle with seasonal data

- ▶ A moving average of the same length as the season removes the seasonal pattern.
- ▶ For quarterly data: use a  $2 \times 4MA$
- ▶ For monthly data: use a  $2 \times 12MA$



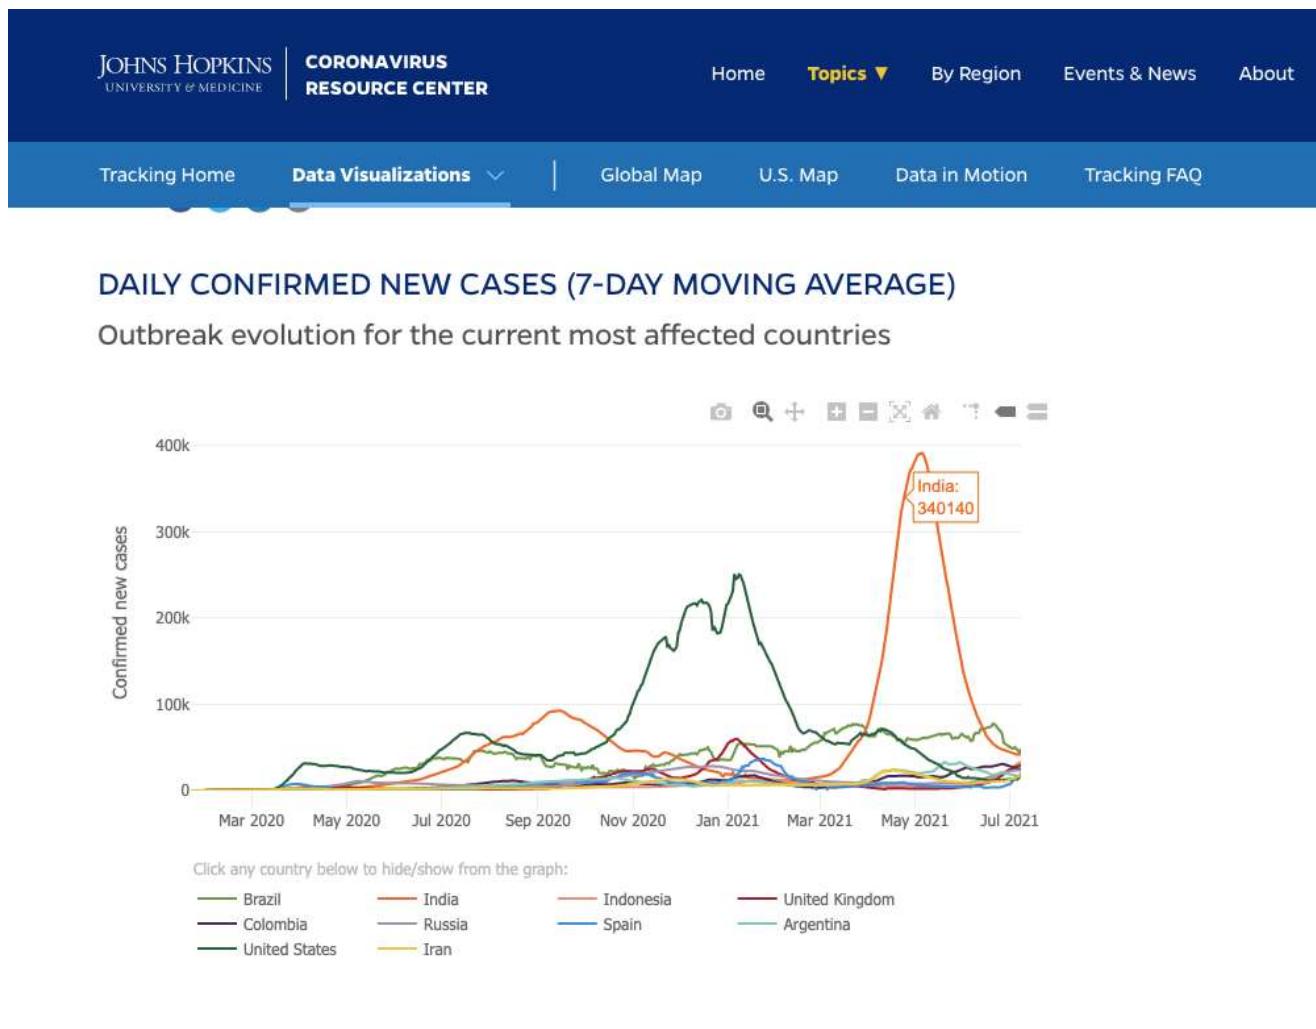
# Weighted moving averages

## Weighted MA

$$T_t = \sum_{j=-k}^k a_j y_{t+j}$$

where  $k = (m - 1)/2$  is the half-width and the weights are denoted by  $[a_{-k}, \dots, a_k]$ .

- ▶ Simple  $m$ -MA: all weights equal to  $1/m$ .
- ▶ Require sum of  $a_j = 1$  and  $a_j = a_{-j}$ .
- ▶ Weighted MA are smoother.



<https://coronavirus.jhu.edu/data/new-cases>

39 / 51

# Trend-cycle

- ▶ Multiplicative decomposition:  $y_t = T_t \times S_t \times R_t = \hat{T}_t \times \hat{S}_t \times \hat{R}_t$
- ▶ Additive decomposition:  $y_t = T_t + S_t + R_t = \hat{T}_t + \hat{S}_t + \hat{R}_t$
- ▶ Estimate  $\hat{T}$  using  $(2 \times m)$ -MA if  $m$  is even. Otherwise, estimate  $\hat{T}$  using  $m$ -MA
- ▶ Compute de-trended series
  - Multiplicative decomposition:  $y_t / \hat{T}_t$
  - Additive decomposition:  $y_t - \hat{T}_t$

## De-trending

Remove smoothed series  $\hat{T}_t$  from  $y_t$  to leave  $S_t$  and  $R_t$ .

- ▶ Multiplicative model:  $\frac{y_t}{\hat{T}_t} = \frac{\hat{T}_t \times \hat{S}_t \times \hat{R}_t}{\hat{T}_t} = \hat{S}_t \times \hat{R}_t$
- ▶ Additive model:  $y_t - \hat{T}_t = (\hat{T}_t + \hat{S}_t + \hat{R}_t) - \hat{T}_t = \hat{S}_t + \hat{R}_t$

# Classical decomposition

- ▶ Choose additive or multiplicative depending on which gives the most stable components.
- ▶ Estimate of trend is unavailable for first few and last few observations.
- ▶ Seasonal component repeats from year to year. May not be realistic.
- ▶ Not robust to outliers.
- ▶ Newer methods designed to overcome these problems.

# History of time series decomposition

# History of time series decomposition

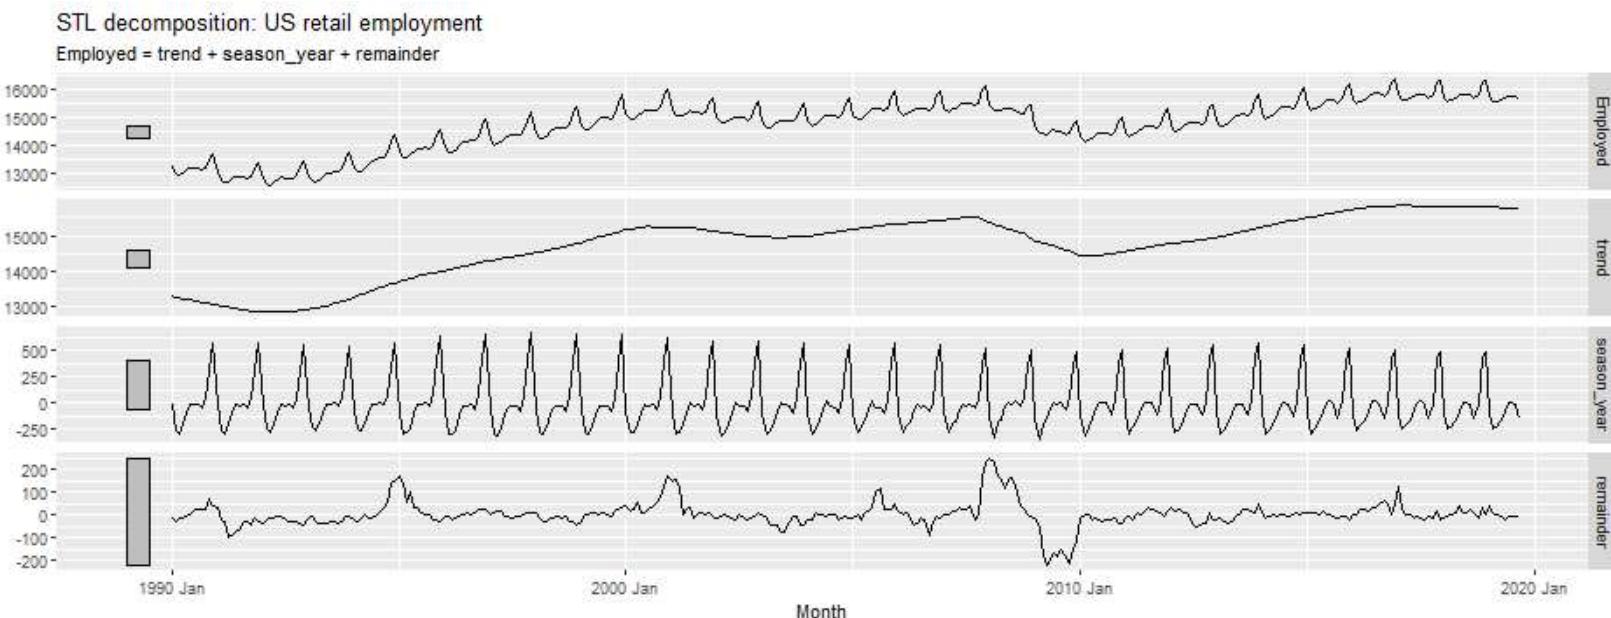
- ▶ Classical method originated in 1920s.
- ▶ Census II method introduced in 1957. Basis for X-11 method and variants (including X-12-ARIMA, X-13-ARIMA)
- ▶ STL method introduced in 1983
- ▶ TRAMO/SEATS introduced in 1990s.

# STL decomposition

- ▶ STL: "Seasonal and Trend decomposition using Loess"
- ▶ Very versatile and robust.
- ▶ Unlike X-12-ARIMA, STL will handle any type of seasonality.
- ▶ Seasonal component allowed to change over time, and rate of change controlled by user.
- ▶ Smoothness of trend-cycle also controlled by user.
- ▶ Robust to outliers
- ▶ Not trading day or calendar adjustments.
- ▶ Only additive.
- ▶ Take logs to get multiplicative decomposition.
- ▶ Use Box-Cox transformations to get other decompositions.

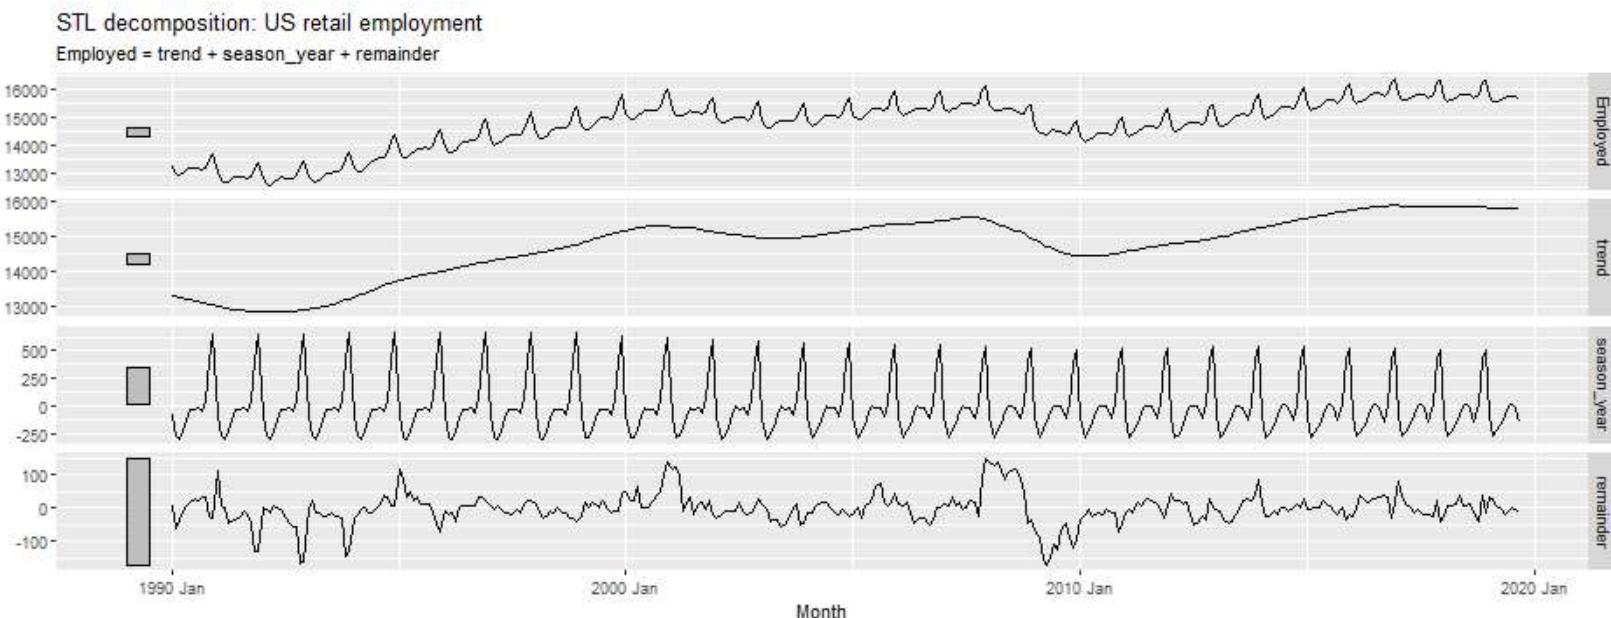
# STL decomposition

```
us_retail_employment |>  
  model(STL(Employed ~ season(window=5), robust=TRUE)) |>  
  components() |> autoplot() +  
  ggtitle("STL decomposition: US retail employment")
```



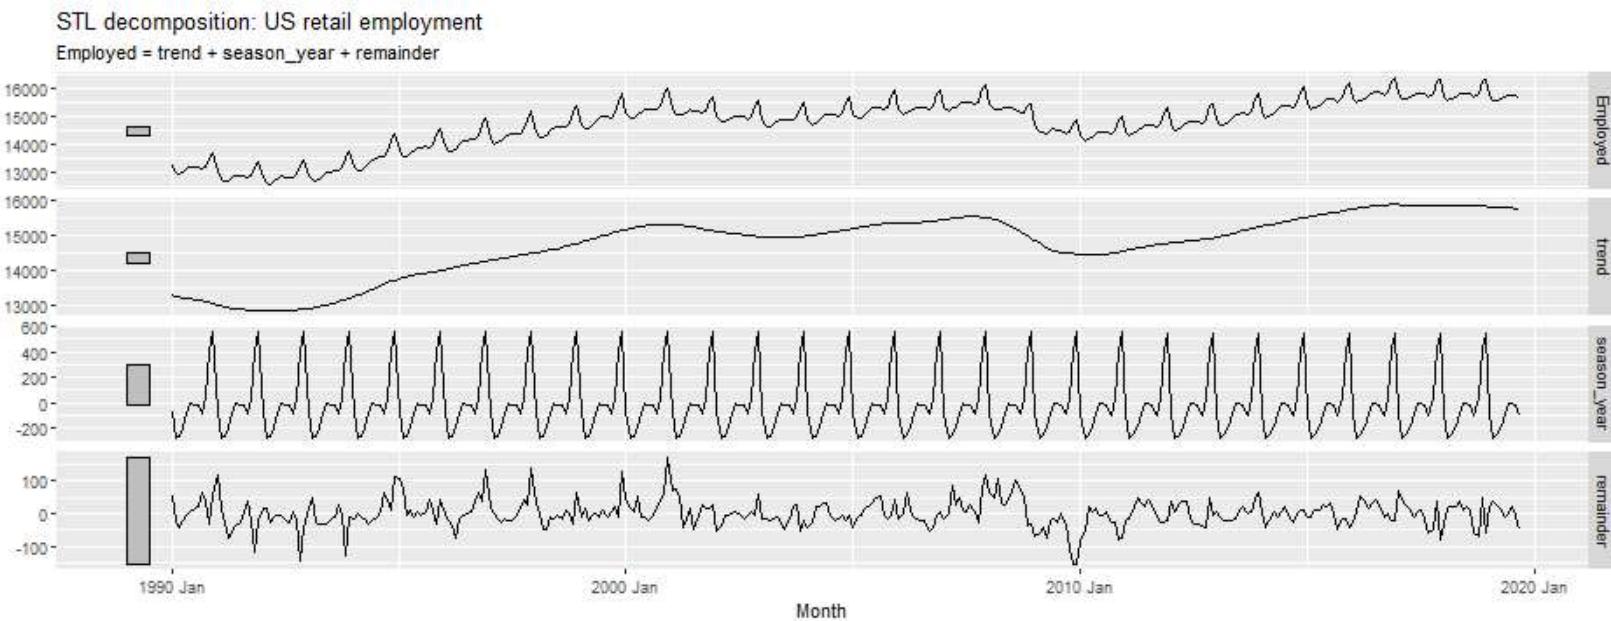
# STL decomposition

```
us_retail_employment |>  
  model(STL(Employed ~ season(window=9), robust=TRUE)) |>  
  components() |> autoplot() +  
  ggtitle("STL decomposition: US retail employment")
```



# STL decomposition

```
us_retail_employment |>  
  model(STL(Employed ~ season(window=55), robust=TRUE)) |>  
  components() |> autoplot() +  
  ggtitle("STL decomposition: US retail employment")
```

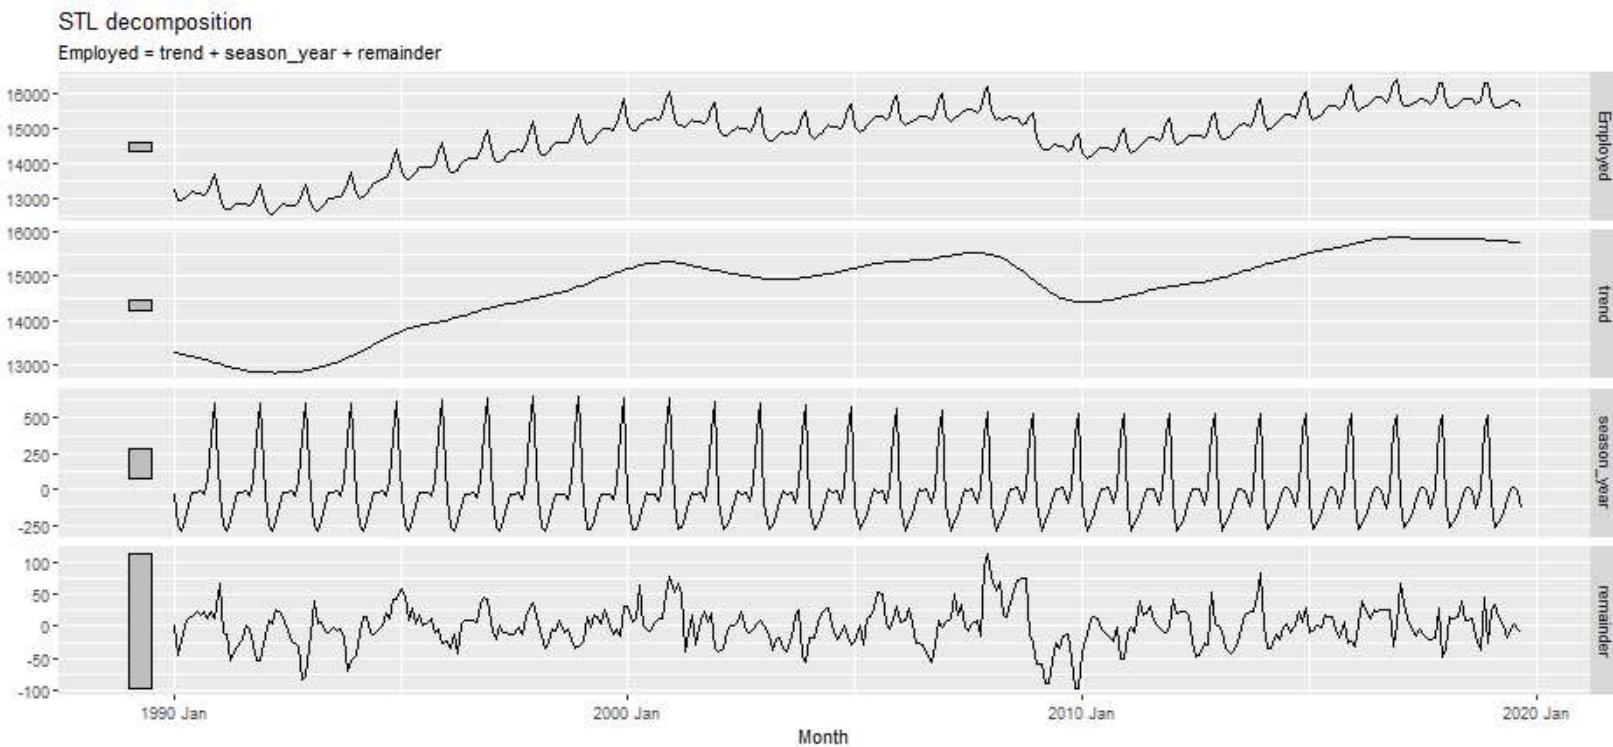


# STL decomposition

```
us_retail_employment |>  
  model(STL(Employed ~ season(window=5))) |>  
  components()  
  
us_retail_employment |>  
  model(STL(Employed ~ trend(window=15) +  
            season(window="periodic"),  
            robust = TRUE)  
) |> components()
```

- ▶ `trend(window = ?)` controls wiggliness of trend component.
- ▶ `season(window = ?)` controls variation on seasonal component.
- ▶ `season(window = 'periodic')` is equivalent to an infinite window.

# STL decomposition



- ▶ `STL()` chooses `season(window=13)` by default
- ▶ Can include transformations.

# STL decomposition

- ▶ Algorithm that updates trend and seasonal components iteratively.
- ▶ Starts with  $\hat{T}_t = 0$
- ▶ Uses a mixture of loess and moving averages to successively refine the trend and seasonal estimates.
- ▶ The trend window controls loess bandwidth applied to deasonalised values.
- ▶ The season window controls loess bandwidth applied to detrended subseries.
- ▶ Default season window = 13
- ▶ Default trend window = `nextodd( ceiling((1.5*period)/(1-(1.5/s.window))) )`

# References

- ▶ Hyndman, R. J., & Athanasopoulos, G. (2018). Forecasting: principles and practice. OTexts.