

# MA 5124 Financial Time Series Analysis and Forecasting

## Chapter 1: Introduction to Time Series & Forecasting Lesson 3

Dr. Priyanga Talagala


Department of Mathematics  
University of Moratuwa

30-06-2024

# **Useful tools for different forecasting situations**


# A tidy forecasting workflow

The process of producing forecasts can be split up into a few fundamental steps.



- 1 Preparing data
- 2 Data visualisation
- 3 Specifying a model
- 4 Model estimation
- 5 Accuracy and performance evaluation
- 6 Producing forecasts


# Some simple forecasting methods



All models are wrong, but some are useful.


— *George E. P. Box* —

# Some simple forecasting methods




How would you forecast these series?

# Some simple forecasting methods



How would you forecast these series?

# Some simple forecasting methods




How would you forecast these series?

# Some simple forecasting methods

## MEAN( $y$ ): Average method

- ▶ Forecasts of all future values is equal to mean of historical data  $\{y_1, \dots, y_T\}$ .
- ▶ Forecasts:  $\hat{y}_{T+h|T} = \bar{y} = (y_1 + \dots + y_T)/T$

```
fit <- bricks |> model(MEAN(Bricks))
```




# Some simple forecasting methods

## NAIVE(y): Naïve method

- ▶ Forecasts equal to last observed value.
- ▶ Forecasts:  $\hat{y}_{T+h|T} = y_T$ .

```
fit <- bricks |> model(NAIVE(Bricks))
```




# Some simple forecasting methods

## SNAIVE( $y \sim \text{lag}(m)$ ): Seasonal naïve method

- ▶ Forecasts equal to last value from same season.
- ▶ Forecasts:  $\hat{y}_{T+h|T} = y_{T+h-m(k+1)}$ , where  $m = \text{seasonal period}$  and  $k$  is the integer part of  $(h - 1)/m$

fit <- bricks |> model(SNAIVE(Bricks~lag("year")))



# Some simple forecasting methods

## RW(y ~ drift()): Drift method

- ▶ Forecasts equal to last value plus average change.
- ▶ Forecasts:

$$\begin{aligned}\hat{y}_{T+h|T} &= y_T + \frac{h}{T-1} \sum_{t=2}^T (y_t - y_{t-1}) \\ &= y_T + \frac{h}{T-1} (y_T - y_1)\end{aligned}$$


.

- ▶ Equivalent to extrapolating a line drawn between first and last observations.

# Some simple forecasting methods

## Drift method

```
bricks |> model(RW(Bricks ~ drift()))
```



# Model fitting

The `model()` function trains models to data.

```
brick_fit <- aus_production |>
  filter(!is.na(Bricks)) |>
  model(
    Seasonal_naive = SNAIVE(Bricks),
    Naive = NAIVE(Bricks),
    Drift = RW(Bricks ~ drift()),
    Mean = MEAN(Bricks)
  )

brick_fit

## # A mable: 1 x 4
##   Seasonal_naive     Naive       Drift      Mean
##   <model> <model>     <model> <model>
## 1 <SNAIVE> <NAIVE> <RW w/ drift> <MEAN>
```

A `mable` is a model table, each cell corresponds to a fitted model.

# Producing forecasts


```
brick_fc <- brick_fit |>  
  forecast(h = "5 years")  
  
print(brick_fc)
```

```
## # A fable: 80 x 4 [1Q]  
## # Key:   .model [4]  
##   .model      Quarter     Bricks .mean  
##   <chr>       <qtr>     <dist> <dbl>  
## 1 Seasonal_naive 2005 Q3 N(428, 2336) 428  
## 2 Seasonal_naive 2005 Q4 N(397, 2336) 397  
## 3 Seasonal_naive 2006 Q1 N(355, 2336) 355  
## 4 Seasonal_naive 2006 Q2 N(435, 2336) 435  
## 5 Seasonal_naive 2006 Q3 N(428, 4672) 428  
## 6 Seasonal_naive 2006 Q4 N(397, 4672) 397  
## 7 Seasonal_naive 2007 Q1 N(355, 4672) 355  
## 8 Seasonal_naive 2007 Q2 N(435, 4672) 435  
## 9 Seasonal_naive 2007 Q3 N(428, 7008) 428  
## 10 Seasonal_naive 2007 Q4 N(397, 7008) 397  
## # i 70 more rows
```

A **fable** is a forecast table with point forecasts and distributions.

# Visualising forecasts

```
brick_fc |>  
  autoplot(aus_production, level = NULL) +  
  ggtitle("Forecasts for quarterly clay brick production") +  
  xlab("Year") + ylab("Millions of bricks") +  
  guides(colour = guide_legend(title = "Forecast"))
```



# Residual diagnostics

# Fitted values

- ▶  $\hat{y}_{t|t-1}$  is the forecast of  $y_t$  based on observations  $y_1, \dots, y_{t-1}$ .
- ▶ We call these "fitted values".
- ▶ Sometimes drop the subscript:  $\hat{y}_t \equiv \hat{y}_{t|t-1}$ .

## For example:

- ▶  $\hat{y}_t = \bar{y}$  for average method.
- ▶  $\hat{y}_t = y_{t-1} + (y_T - y_1)/(T - 1)$  for drift method.

# Forecasting residuals

**Residuals in forecasting:** difference between observed value and its fitted value:  $e_t = y_t - \hat{y}_{t|t-1}$ .

## Assumptions


- 1  $\{e_t\}$  uncorrelated. If they aren't, then information left in residuals that should be used in computing forecasts.
- 2  $\{e_t\}$  have mean zero. If they don't, then forecasts are biased.

## Useful properties (for distributions and prediction intervals)

- 3  $\{e_t\}$  have constant variance.
- 4  $\{e_t\}$  are normally distributed.

# Facebook closing stock price

```
fb_stock <- gafa_stock |>  
  filter(Symbol == "FB") |>  
  mutate(trading_day = row_number()) |>  
  update_tsibble(index = trading_day, regular = TRUE)  
fb_stock |> autoplot(Close)
```



# Facebook closing stock price

```
fit <- fb_stock |> model(NAIVE(Close))  
fit
```


```
# A mable: 1 x 2  
# Key:     Symbol [1]  
Symbol `NAIVE(Close)`  
<chr>      <model>  
1 FB        <NAIVE>
```

```
augment(fit)
```

```
# A tsibble: 1,258 x 7 [1]  
# Key:     Symbol, .model [1]  
Symbol .model      trading_day Close .fitted .resid .innov  
<chr>  <chr>          <int>  <dbl>   <dbl>   <dbl>   <dbl>  
1 FB    NAIVE(Close)       1  54.7    NA    NA    NA  
2 FB    NAIVE(Close)       2  54.6  54.7 -0.150 -0.150  
3 FB    NAIVE(Close)       3  57.2  54.6  2.64  2.64  
4 FB    NAIVE(Close)       4  57.9  57.2  0.720 0.720  
5 FB    NAIVE(Close)       5  58.2  57.9  0.310 0.310  
6 FB    NAIVE(Close)       6  57.2  58.2 -1.01 -1.01  
7 FB    NAIVE(Close)       7  57.9  57.2  0.720 0.720  
8 FB    NAIVE(Close)      8  55.9  57.9 -2.03 -2.03
```


# Facebook closing stock price

```
augment(fit) |>  
  ggplot(aes(x = trading_day)) +  
  geom_line(aes(y = Close, colour = "Data")) +  
  geom_line(aes(y = .fitted, colour = "Fitted"))
```




# Facebook closing stock price

```
augment(fit) |>  
  filter(trading_day > 1100) |>  
  ggplot(aes(x = trading_day)) +  
  geom_line(aes(y = Close, colour = "Data")) +  
  geom_line(aes(y = .fitted, colour = "Fitted"))
```




# Facebook closing stock price

```
augment(fit) |>  
  autoplot(.resid) + xlab("Day") + ylab("") +  
  ggtitle("Residuals from naïve method")
```




# Facebook closing stock price

```
augment(fit) |>  
  ggplot(aes(x = .resid)) +  
  geom_histogram(bins = 150) +  
  ggtitle("Histogram of residuals")
```



# Facebook closing stock price

```
augment(fit) |>  
  ACF(.resid) |>  
  autoplot() + ggtitle("ACF of residuals")
```




## ACF of residuals

- ▶ We assume that the residuals are white noise (uncorrelated, mean zero, constant variance). If they aren't, then there is information left in the residuals that should be used in computing forecasts.
- ▶ So a standard residual diagnostic is to check the ACF of the residuals of a forecasting method.
- ▶ We *expect* these to look like white noise.

# gg\_tsresiduals function

```
gg_tsresiduals(fit)
```



# Portmanteau tests

Consider a *whole set* of  $r_k$  values, and develop a test to see whether the set is significantly different from a zero set.

## Box-Pierce test

$$Q = T \sum_{k=1}^h r_k^2$$

where  $h$  is max lag being considered and  $T$  is number of observations.

- ▶ If each  $r_k$  close to zero,  $Q$  will be **small**.
- ▶ If some  $r_k$  values large (positive or negative),  $Q$  will be **large**.

# Portmanteau tests

Consider a *whole set* of  $r_k$  values, and develop a test to see whether the set is significantly different from a zero set.

## Ljung-Box test

$$Q^* = T(T + 2) \sum_{k=1}^h (T - k)^{-1} r_k^2$$

where  $h$  is max lag being considered and  $T$  is number of observations.

- ▶ My preferences:  $h = 10$  for non-seasonal data,  $h = 2m$  for seasonal data.
- ▶ Better performance, especially in small samples.

# Portmanteau tests

- ▶ If data are WN,  $Q^*$  has  $\chi^2$  distribution with  $(h - K)$  degrees of freedom where  $K = \text{no. parameters in model}$ .
- ▶ When applied to raw data, set  $K = 0$ .

```
augment(fit) |>
  features(.resid, ljung_box, lag=10, dof=0)

## # A tibble: 1 × 4
##   Symbol .model      lb_stat lb_pvalue
##   <chr>  <chr>       <dbl>     <dbl>
## 1 FB     NAIVE(Close) 12.1      0.276
```

# Distributional forecasts and prediction intervals

# Forecast distributions

- ▶ A forecast  $\hat{y}_{T+h|T}$  is (usually) the mean of the conditional distribution  $y_{T+h} \mid y_1, \dots, y_T$ .
- ▶ Most time series models produce normally distributed forecasts.
- ▶ The forecast distribution describes the probability of observing any future value.

# Forecast distributions

Assuming residuals are normal, uncorrelated,  $\text{sd} = \hat{\sigma}$ :

- ▶ **Mean:**  $\hat{y}_{T+h|T} \sim N(\bar{y}, (1 + 1/T)\hat{\sigma}^2)$
- ▶ **Naïve:**  $\hat{y}_{T+h|T} \sim N(y_T, h\hat{\sigma}^2)$
- ▶ **Seasonal naïve**  $\hat{y}_{T+h|T} \sim N(y_{T+h-m(k+1)}, (k+1)\hat{\sigma}^2)$
- ▶ **Drift:**  $\hat{y}_{T+h|T} \sim N(y_T + \frac{h}{T-1}(y_T - y_1), h\frac{T+h}{T}\hat{\sigma}^2)$

where  $k$  is the integer part of  $(h - 1)/m$ .

Note that when  $h = 1$  and  $T$  is large, these all give the same approximate forecast variance:  $\hat{\sigma}^2$ .

## Prediction intervals

- ▶ A prediction interval gives a region within which we expect  $y_{T+h}$  to lie with a specified probability.
- ▶ Assuming forecast errors are normally distributed, then a 95% PI is

$$\hat{y}_{T+h|T} \pm 1.96\hat{\sigma}_h$$

where  $\hat{\sigma}_h$  is the st dev of the  $h$ -step distribution.


- ▶ When  $h = 1$ ,  $\hat{\sigma}_h$  can be estimated from the residuals.

# Prediction intervals

```
brick_fc |> hilo(level = 95)
```

```
## # A tsibble: 80 x 5 [1Q]
## # Key:      .model [4]
##   .model      Quarter     Bricks .mean      `95%` 
##   <chr>       <qtr>      <dist> <dbl>      <hilo>
## 1 Seasonal_naive 2005 Q3 N(428, 2336)  428 [333.2737, 522.7263]95
## 2 Seasonal_naive 2005 Q4 N(397, 2336)  397 [302.2737, 491.7263]95
## 3 Seasonal_naive 2006 Q1 N(355, 2336)  355 [260.2737, 449.7263]95
## 4 Seasonal_naive 2006 Q2 N(435, 2336)  435 [340.2737, 529.7263]95
## 5 Seasonal_naive 2006 Q3 N(428, 4672)  428 [294.0368, 561.9632]95
## 6 Seasonal_naive 2006 Q4 N(397, 4672)  397 [263.0368, 530.9632]95
## 7 Seasonal_naive 2007 Q1 N(355, 4672)  355 [221.0368, 488.9632]95
## 8 Seasonal_naive 2007 Q2 N(435, 4672)  435 [301.0368, 568.9632]95
## 9 Seasonal_naive 2007 Q3 N(428, 7008)  428 [263.9292, 592.0708]95
## 10 Seasonal_naive 2007 Q4 N(397, 7008) 397 [232.9292, 561.0708]95
## # i 70 more rows
```

```
bricks |>
  filter(!is.na(Bricks)) |>
  model(
    Seasonal_naive = SNAIVE(Bricks),
    Naive = NAIVE(Bricks),
    Drift = RW(Bricks ~ drift()),
    Mean = MEAN(Bricks)
  ) |>
  forecast(h = 10) |> autoplot(bricks)
```



# Prediction intervals

- ▶ Point forecasts are often useless without a measure of uncertainty (such as prediction intervals).
- ▶ Prediction intervals require a stochastic model (with random errors, etc).
- ▶ Multi-step forecasts for time series require a more sophisticated approach (with PI getting wider as the forecast horizon increases).

# Prediction intervals

- ▶ Computed automatically from the forecast distribution.
- ▶ Use `level` argument to control coverage.
- ▶ Check residual assumptions before believing them (we will see this next class).
- ▶ Usually too narrow due to unaccounted uncertainty.

# Forecasting and decomposition

# Forecasting and decomposition

- ▶ Forecast seasonal component by repeating the last year
- ▶ Forecast seasonally adjusted data using non-seasonal time series method.
- ▶ Combine forecasts of seasonal component with forecasts of seasonally adjusted data to get forecasts of original data.
- ▶ Sometimes a decomposition is useful just for understanding the data before building a separate forecasting model.

# US Retail Employment

```
us_retail_employment <- us_employment |>
  filter(year(Month) >= 1990, Title == "Retail Trade") |>
  select(-Series_ID)
us_retail_employment
```

```
## # A tsibble: 357 x 3 [1M]
##   Month Title     Employed
##   <mth> <chr>     <dbl>
## 1 1990 Jan Retail Trade 13256.
## 2 1990 Feb Retail Trade 12966.
## 3 1990 Mar Retail Trade 12938.
## 4 1990 Apr Retail Trade 13012.
## 5 1990 May Retail Trade 13108.
## 6 1990 Jun Retail Trade 13183.
## 7 1990 Jul Retail Trade 13170.
## 8 1990 Aug Retail Trade 13160.
## 9 1990 Sep Retail Trade 13113.
## 10 1990 Oct Retail Trade 13185.
## # i 347 more rows
```


# US Retail Employment

```
dcmp <- us_retail_employment |>
  model(STL(Employed)) |>
  components() |>
  select(-.model)
dcmp

## # A tsibble: 357 x 6 [1M]
##   Month Employed trend season_year remainder season_adjust
##   <mth>    <dbl>  <dbl>     <dbl>    <dbl>      <dbl>
## 1 1990 Jan    13256. 13288.    -33.0    0.836    13289.
## 2 1990 Feb    12966. 13269.    -258.    -44.6    13224.
## 3 1990 Mar    12938. 13250.    -290.    -22.1    13228.
## 4 1990 Apr    13012. 13231.    -220.     1.05    13232.
## 5 1990 May    13108. 13211.    -114.     11.3    13223.
## 6 1990 Jun    13183. 13192.    -24.3     15.5    13207.
## 7 1990 Jul    13170. 13172.    -23.2     21.6    13193.
## 8 1990 Aug    13160. 13151.    -9.52    17.8    13169.
## 9 1990 Sep    13113. 13131.   -39.5    22.0    13153.
## 10 1990 Oct   13185. 13110.    61.6     13.2    13124.
## # i 347 more rows
```


# US Retail Employment

```
dcmp |>  
  model(NAIVE(season_adjust)) |>  
  forecast() |>  
  autoplot(dcmp) +  
  ggtitle("Naive forecasts of seasonally adjusted data")
```



# US Retail Employment

```
fit_dcmp <- us_retail_employment |>  
  model(stlf = decomposition_model(  
    STL(Employed ~ trend(window = 7), robust = TRUE),  
    NAIVE(season_adjust)  
)  
  
fit_dcmp |> forecast() |>  
  autoplot(us_retail_employment)
```




# Decomposition models

`decomposition_model()` creates a decomposition model


- ▶ You must provide a method for forecasting the `season_adjust` series.
- ▶ A seasonal naive method is used by default for the `seasonal` components.
- ▶ The variances from both the seasonally adjusted and seasonal forecasts are combined.

```
fit_dcmp |> gg_tsresiduals()
```



# Evaluating forecast accuracy

# Training and test sets



- ▶ A model which fits the training data well will not necessarily forecast well.
  - ▶ A perfect fit can always be obtained by using a model with enough parameters.
  - ▶ Over-fitting a model to data is just as bad as failing to identify a systematic pattern in the data.
  - ▶ The test set must not be used for *any* aspect of model development or calculation of forecasts.
  - ▶ Forecast accuracy is based only on the test set.
-

# Forecast errors


Forecast "error": the difference between an observed value and its forecast.

$$e_{T+h} = y_{T+h} - \hat{y}_{T+h|T},$$

where the training data is given by  $\{y_1, \dots, y_T\}$

- ▶ Unlike residuals, forecast errors on the test set involve multi-step forecasts.
- ▶ These are *true* forecast errors as the test data is not used in computing  $\hat{y}_{T+h|T}$ .

# Measures of forecast accuracy



# Measures of forecast accuracy

$y_{T+h}$  =  $(T + h)$ th observation,  $h = 1, \dots, H$

$\hat{y}_{T+h|T}$  = its forecast based on data up to time  $T$ .

$$e_{T+h} = y_{T+h} - \hat{y}_{T+h|T}$$

$$\text{MAE} = \text{mean}(|e_{T+h}|)$$

$$\text{MSE} = \text{mean}(e_{T+h}^2)$$

$$\text{RMSE} = \sqrt{\text{mean}(e_{T+h}^2)}$$

$$\text{MAPE} = 100\text{mean}(|e_{T+h}|/|y_{T+h}|)$$

- ▶ MAE, MSE, RMSE are all scale dependent.
- ▶ MAPE is scale independent but is only sensible if  $y_t \gg 0$  for all  $t$ , and  $y$  has a natural zero.

# Measures of forecast accuracy

## Mean Absolute Scaled Error

$$\text{MASE} = \text{mean}(|e_{T+h}|)/Q)$$

where  $Q$  is a stable measure of the scale of the time series  $\{y_t\}$ .

Proposed by Hyndman and Koehler (IJF, 2006).

For non-seasonal time series,

$$Q = \frac{1}{T-1} \sum_{t=2}^T |y_t - y_{t-1}|$$

works well. Then MASE is equivalent to MAE relative to a naïve method.

# Measures of forecast accuracy

## Mean Absolute Scaled Error

$$\text{MASE} = \text{mean}(|e_{T+h}|)/Q)$$

where  $Q$  is a stable measure of the scale of the time series  $\{y_t\}$ .


Proposed by Hyndman and Koehler (IJF, 2006).

For seasonal time series,

$$Q = \frac{1}{T-m} \sum_{t=m+1}^T |y_t - y_{t-m}|$$

works well. Then MASE is equivalent to MAE relative to a seasonal naïve method.

# Measures of forecast accuracy



# Measures of forecast accuracy

```
recent_production <- aus_production |>
  filter(year(Quarter) >= 1992)
train <- recent_production |>
  filter(year(Quarter) <= 2007)
beer_fit <- train |>
  model(
    Mean = MEAN(Beer),
    Naive = NAIVE(Beer),
    Seasonal_naive = SNAIVE(Beer),
    Drift = RW(Beer ~ drift())
  )
beer_fc <- beer_fit |>
  forecast(h = 10)
```

# Measures of forecast accuracy

```
accuracy(beer_fit)
```

```
## # A tibble: 4 × 6
##   .model      .type    RMSE    MAE    MAPE    MASE
##   <chr>       <chr>    <dbl>   <dbl>   <dbl>   <dbl>
## 1 Drift       Training  65.3   54.8   12.2   3.83
## 2 Mean        Training  43.6   35.2   7.89   2.46
## 3 Naive       Training  65.3   54.7   12.2   3.83
## 4 Seasonal_naive Training 16.8   14.3   3.31   1
```

```
accuracy(beer_fc, recent_production)
```

```
## # A tibble: 4 × 6
##   .model      .type    RMSE    MAE    MAPE    MASE
##   <chr>       <chr>    <dbl>   <dbl>   <dbl>   <dbl>
## 1 Drift       Test     64.9   58.9   14.6   4.12
## 2 Mean        Test     38.4   34.8   8.28   2.44
## 3 Naive       Test     62.7   57.4   14.2   4.01
## 4 Seasonal_naive Test    14.3   13.4   3.17   0.937
```

## Poll: true or false?

- 1 Good forecast methods should have normally distributed residuals.
- 2 A model with small residuals will give good forecasts.
- 3 The best measure of forecast accuracy is MAPE.
- 4 If your model doesn't forecast well, you should make it more complicated.
- 5 Always choose the model with the best forecast accuracy as measured on the test set.


# Time series cross-validation

# Time series cross-validation

## Traditional evaluation



## Time series cross-validation



- ▶ Forecast accuracy averaged over test sets.
- ▶ Also known as "evaluation on a rolling forecasting origin"

# Creating the rolling training sets

There are three main rolling types which can be used.

- ▶ Stretch: extends a growing length window with new data.
- ▶ Slide: shifts a fixed length window through the data.
- ▶ Tile: moves a fixed length window without overlap.

Three functions to roll a tsibble: `stretch_tsibble()`, `slide_tsibble()`, and `tile_tsibble()`.

For time series cross-validation, stretching windows are most commonly used.

A good way to choose the best forecasting model is to find the model with the smallest RMSE computed using time series cross-validation.

# Time series cross-validation

Stretch with a minimum length of 3, growing by 1 each step.

```
fb_stretch <- fb_stock |>  
  stretch_tsibble(.init = 3, .step = 1) |>  
  filter(.id != max(.id))
```

```
## # A tsibble: 790,650 x 4 [1]  
## # Key:      .id [1,255]  
##   Date      Close trading_day   .id  
##   <date>    <dbl>     <int> <int>  
## 1 2014-01-02  54.7        1     1  
## 2 2014-01-03  54.6        2     1  
## 3 2014-01-06  57.2        3     1  
## 4 2014-01-02  54.7        1     2  
## 5 2014-01-03  54.6        2     2  
## 6 2014-01-06  57.2        3     2  
## 7 2014-01-07  57.9        4     2  
## 8 2014-01-02  54.7        1     3  
## 9 2014-01-03  54.6        2     3  
## 10 2014-01-06  57.2        3     3  
## 11 2014-01-07  57.9        4     3  
## 12 2014-01-08  58.2        5     3  
## 13 2014-01-02  54.7        1     4
```

# Time series cross-validation

Estimate RW w/ drift models for each window.

```
fit_cv <- fb_stretch |>
  model(RW(Close ~ drift()))

## # A mable: 1,255 x 3
## # Key:     .id, Symbol [1,255]
##   .id Symbol `RW(Close ~ drift())` 
##   <int> <chr>           <model>
## 1     1 FB           <RW w/ drift>
## 2     2 FB           <RW w/ drift>
## 3     3 FB           <RW w/ drift>
## 4     4 FB           <RW w/ drift>
## # ... with 1,251 more rows
```

# Time series cross-validation

Produce one step ahead forecasts from all models.

```
fc_cv <- fit_cv |>
  forecast(h=1)

## # A fable: 1,255 x 5 [1]
## # Key:     .id, Symbol [1,255]
##   .id Symbol trading_day     Close .mean
##   <int> <chr>        <dbl>     <dist> <dbl>
## 1     1 FB            4 N(58, 5.8) 58.4
## 2     2 FB            5 N(59, 2.7) 59.0
## 3     3 FB            6 N(59, 1.9) 59.1
## 4     4 FB            7 N(58, 2.2) 57.7
## # i 1,251 more rows
```

# Time series cross-validation

```
# Time series cross-validation accuracy  
fc_cv |> accuracy(fb_stock)  
# Training set (Residual accuracy)  
fb_stock |> model(RW(Close ~ drift())) |> accuracy()
```

	<b>RMSE</b>	<b>MAE</b>	<b>MAPE</b>
Cross-validation	2.418	1.469	1.266
Training	2.414	1.465	1.261

# References

- ▶ Hyndman, R. J., & Athanasopoulos, G. (2018). Forecasting: principles and practice. OTexts.